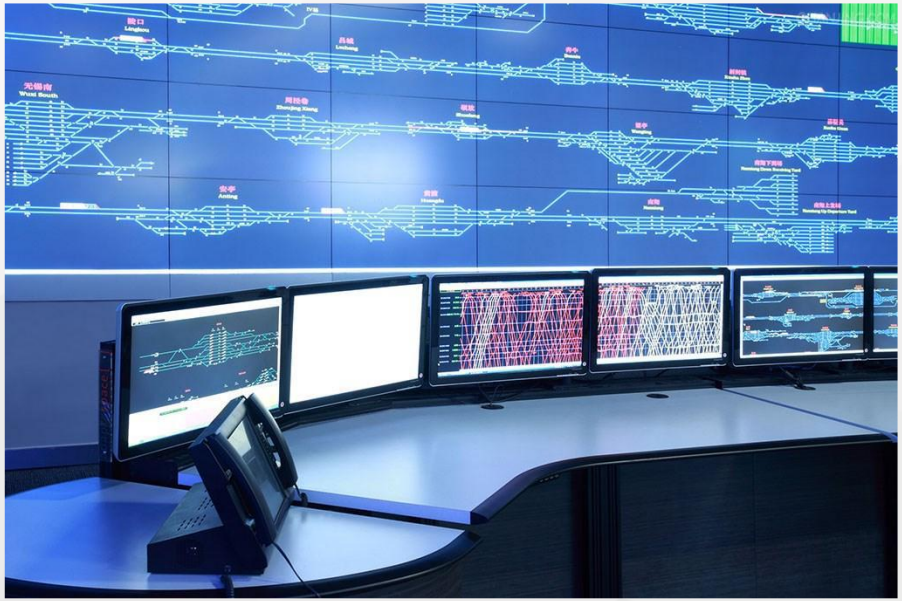

Hadoop Map-Reduce

2017.2 XenRon

CONTENTS

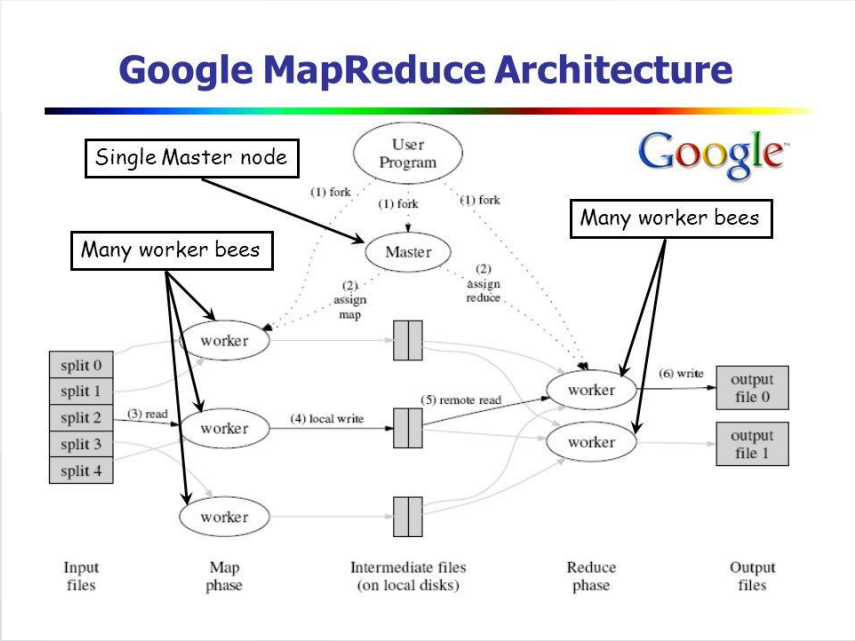

Computation Modal

Divide & Conquer

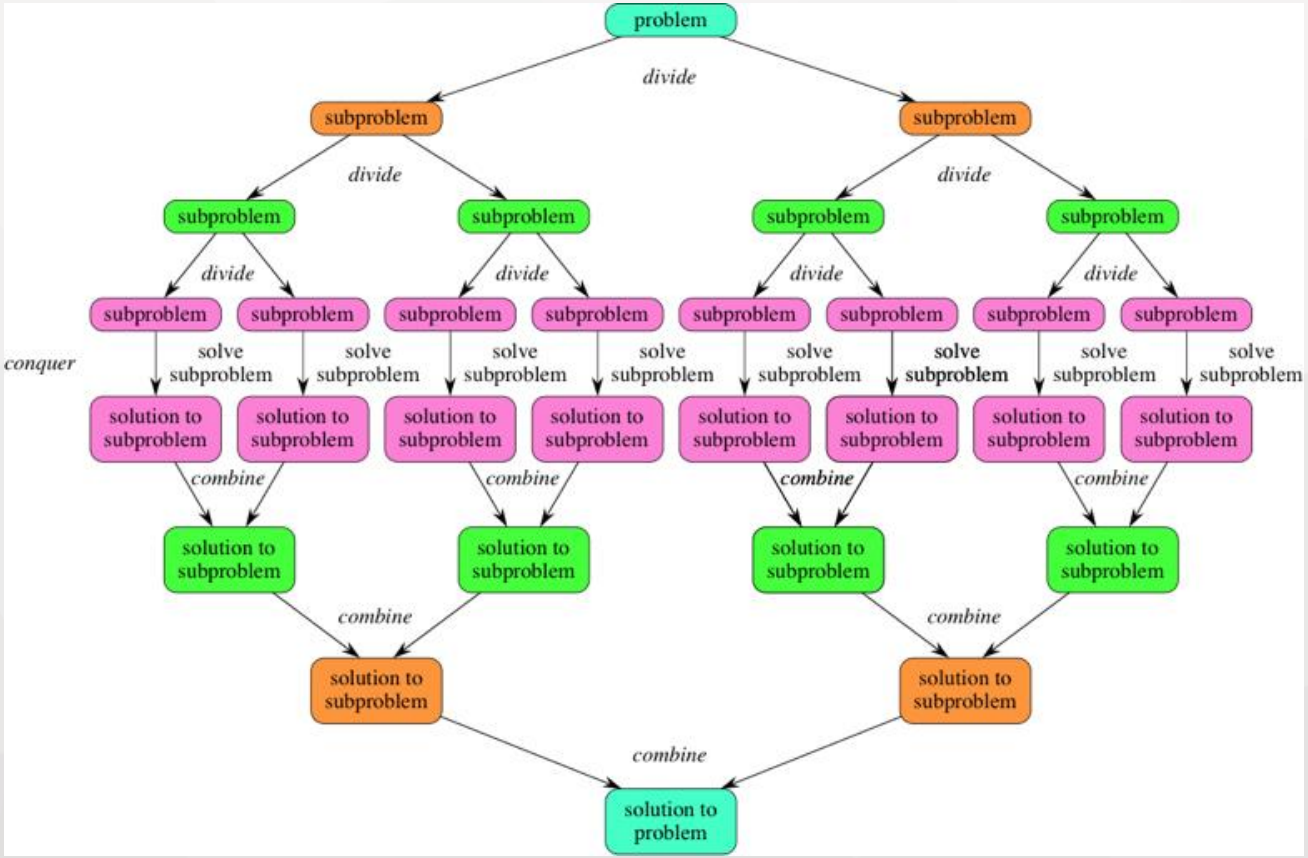

Divide & Conquer

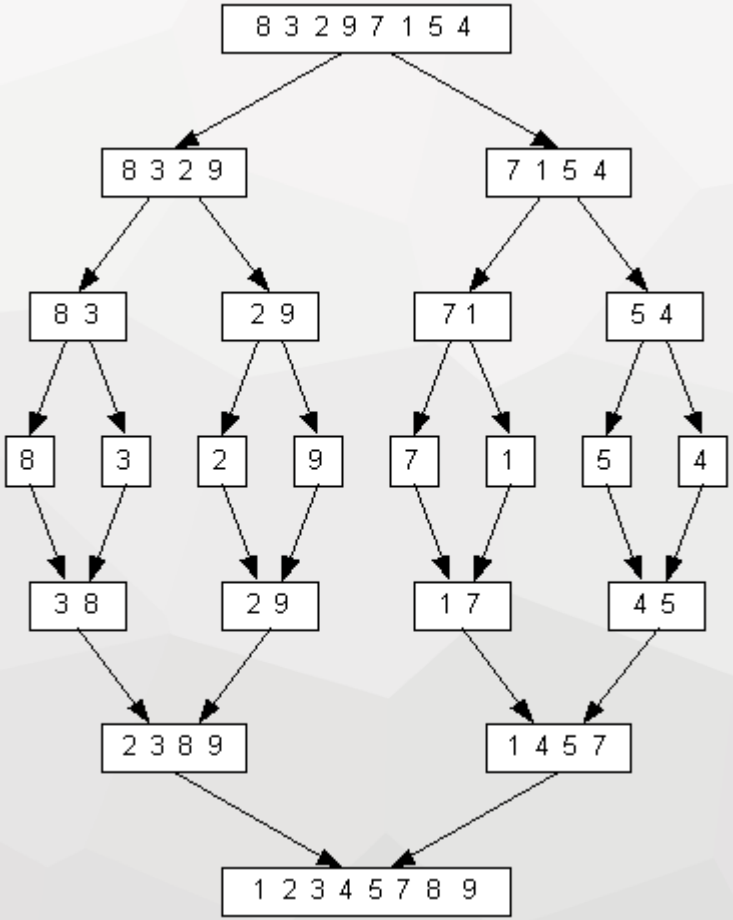

Map-Reduce

Google

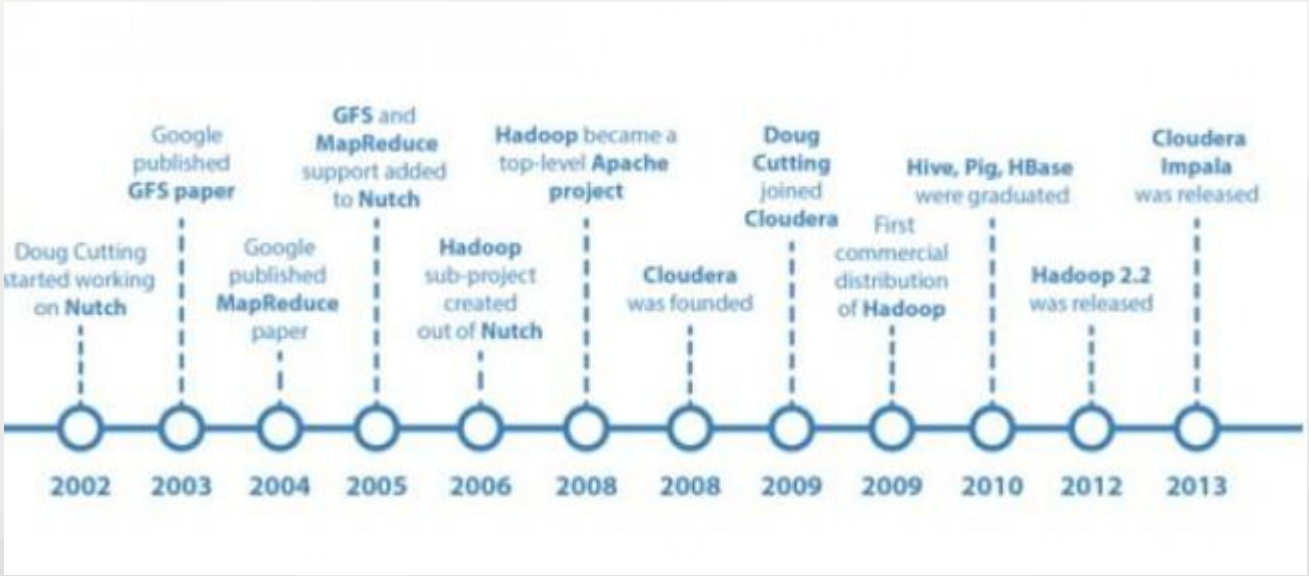
Doug Cutting started working on Nutch
2002
Google published GFS paper
2003
Google published MapReduce paper
2004
GFS and MapReduce support added to Nutch
2005
Hadoop sub-project created out of Nutch
2006
Hadoop became a top-level Apache project
2008
Cloudera was founded
2008
Doug Cutting joined Cloudera
2009
First commercial distribution of Hadoop
2009
Hive, Pig, HBase were graduated
2010
Hadoop 2.2 was released
2012
Cloudera Impala was released
2013

Map Reduce

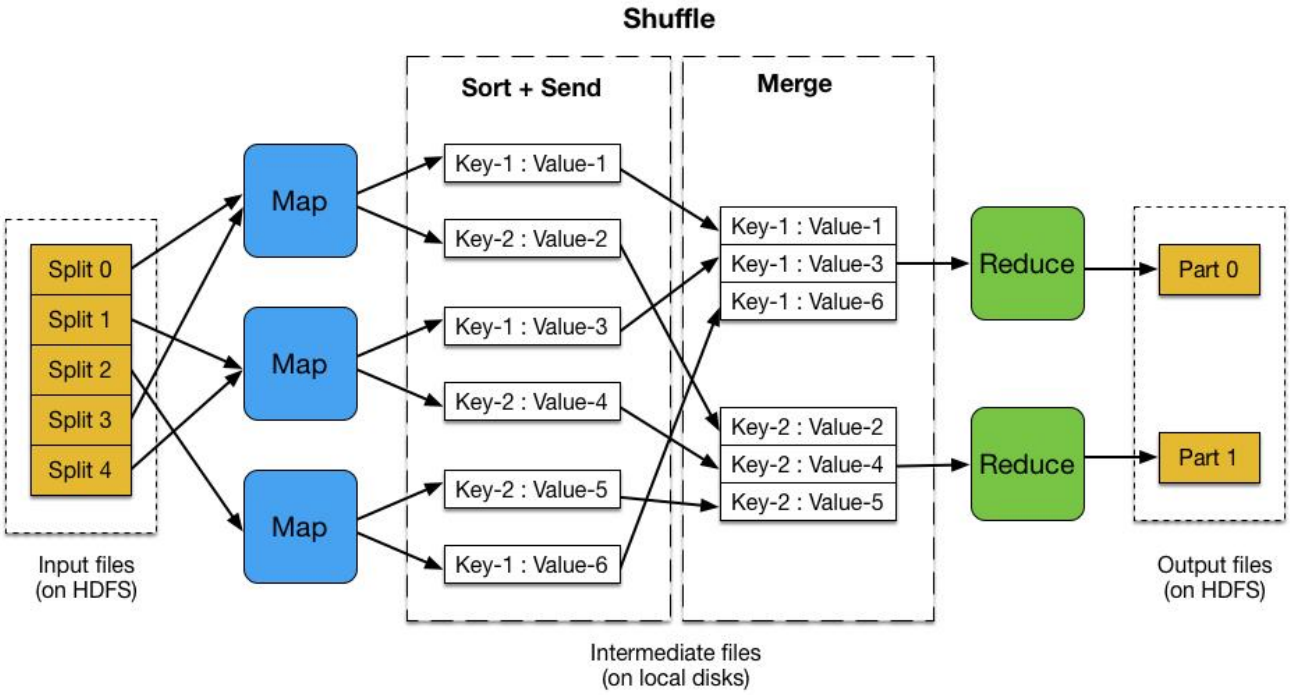
Shuffle
Sort + Send
Merge
Split 0
Split 1
Split 2
Split 3
Split 4
Map
Map
Map
Key-1 : Value-1
Key-2 : Value-2
Key-1 : Value-3
Key-2 : Value-4
Key-2 : Value-5
Key-1 : Value-6
Key-1 : Value-1
Key-1 : Value-3
Key-1 : Value-6
Key-2 : Value-2
Key-2 : Value-4
Key-2 : Value-5
Reduce
Reduce
Part 0
Part 1
Input files
(on HDFS)
Intermediate files
(on local disks)
Output files
(on HDFS)

Map Reduce

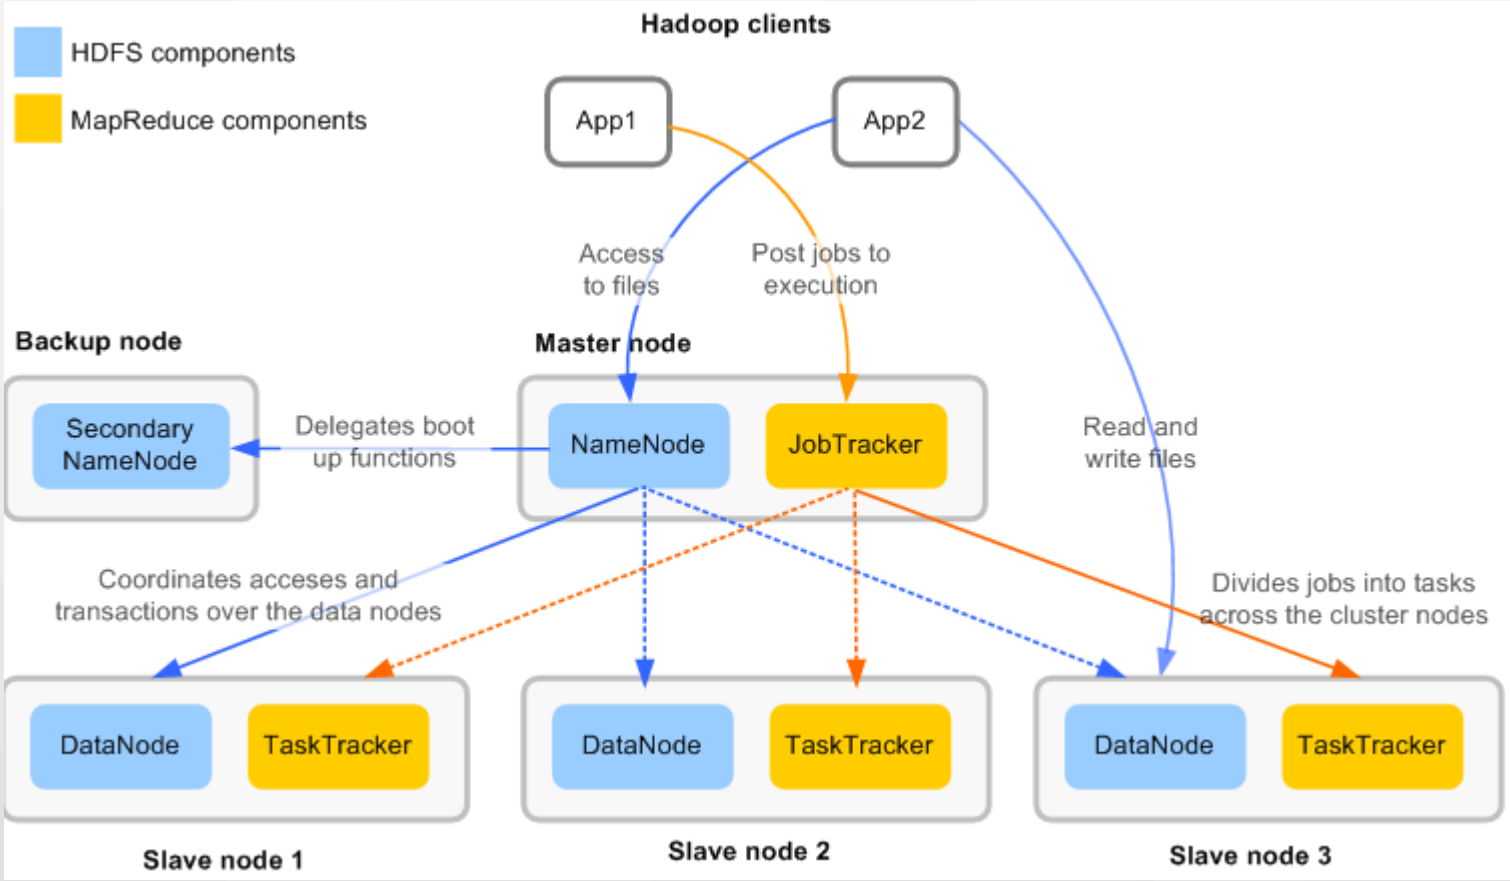
Hadoop clients
HDFS components
MapReduce components
App1
App2
Access to files
Post jobs to execution
Backup node
Master node
Secondary NameNode
Delegates boot up functions
NameNode
JobTracker
Read and write files
Coordinates acceses and transactions over the data nodes
Divides jobs into tasks across the cluster nodes
DataNode
TaskTracker
DataNode
TaskTracker
DataNode
TaskTracker
Slave node 1
Slave node 2
Slave node 3

Map Reduce

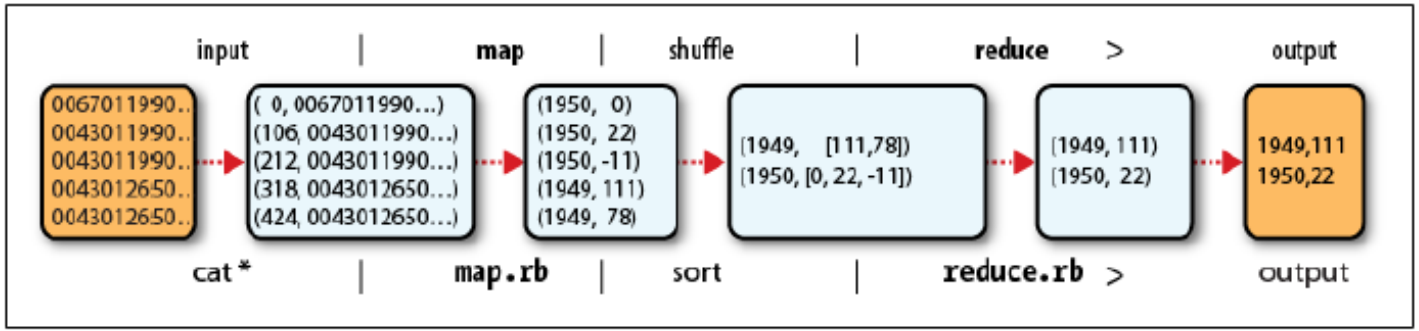

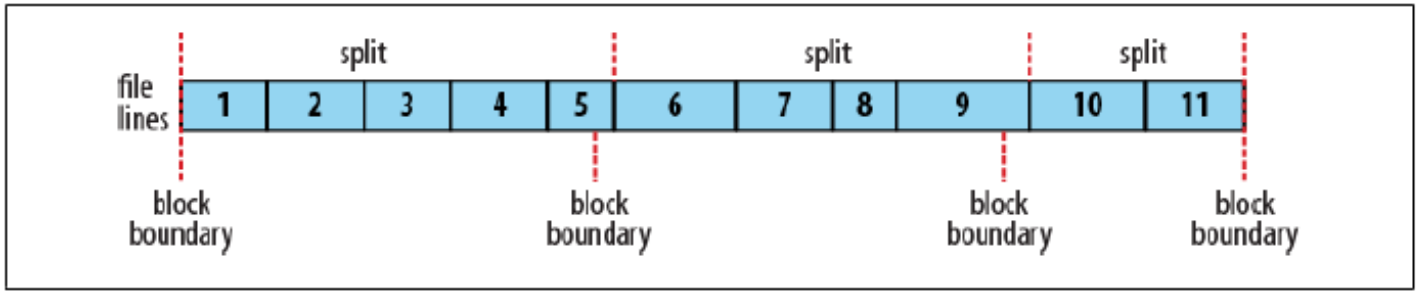

Figure 7-3. Logical records and HDFS blocks for TextInputFormat

Map Reduce

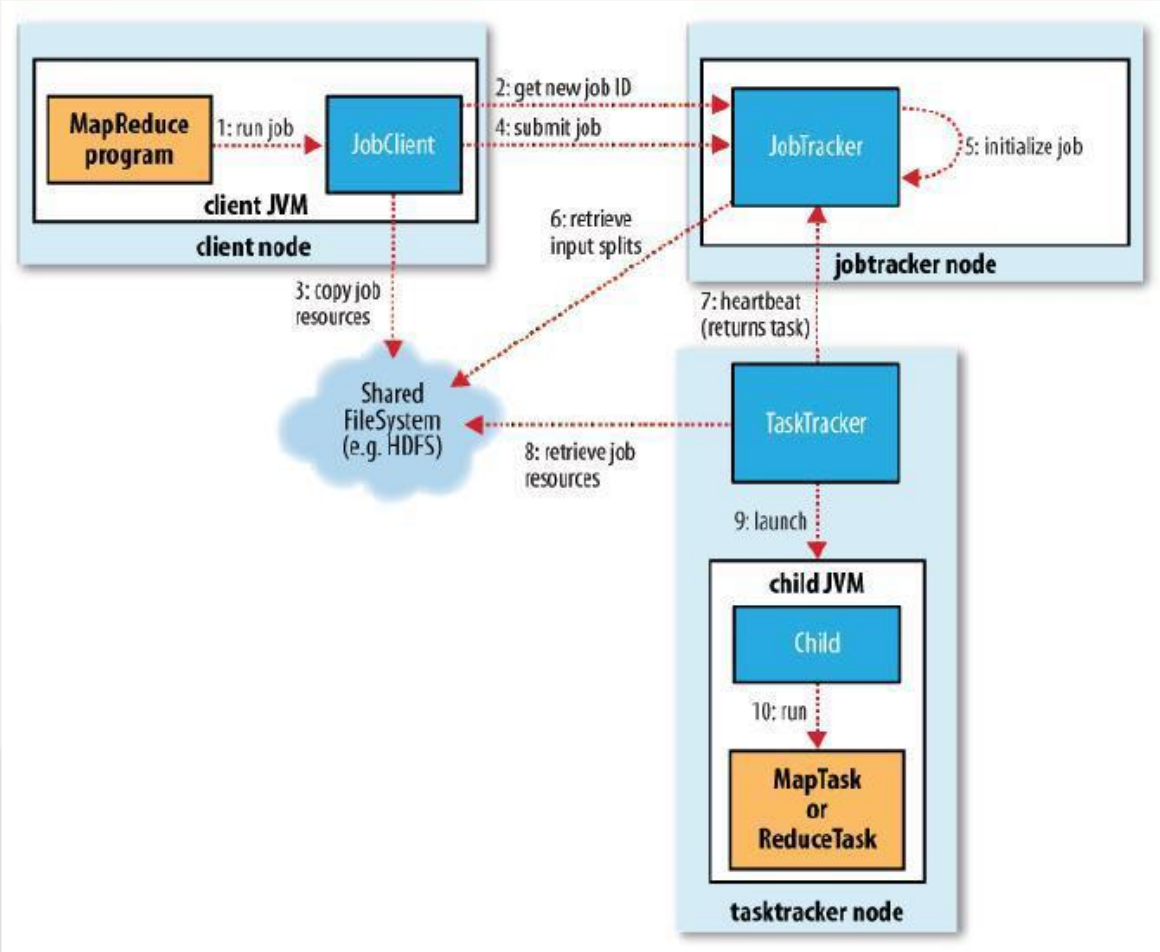
MapReduce program
1: run job
JobClient
client JVM
client node
2: get new job ID
4: submit job
JobTracker
5: initialize job
jobtracker node
6: retrieve input splits
3: copy job resources
7: heartbeat (returns task)
Shared FileSystem (e.g. HDFS)
8: retrieve job resources
TaskTracker
9: launch
child JVM
Child
10: run
MapTask or ReduceTask
tasktracker node

Algorithms in Mahout

See http://cwiki.apache.org/confluence/display/MAHOUT/Algorithms

etu
nexus enlightenment

Word Count

Word Count

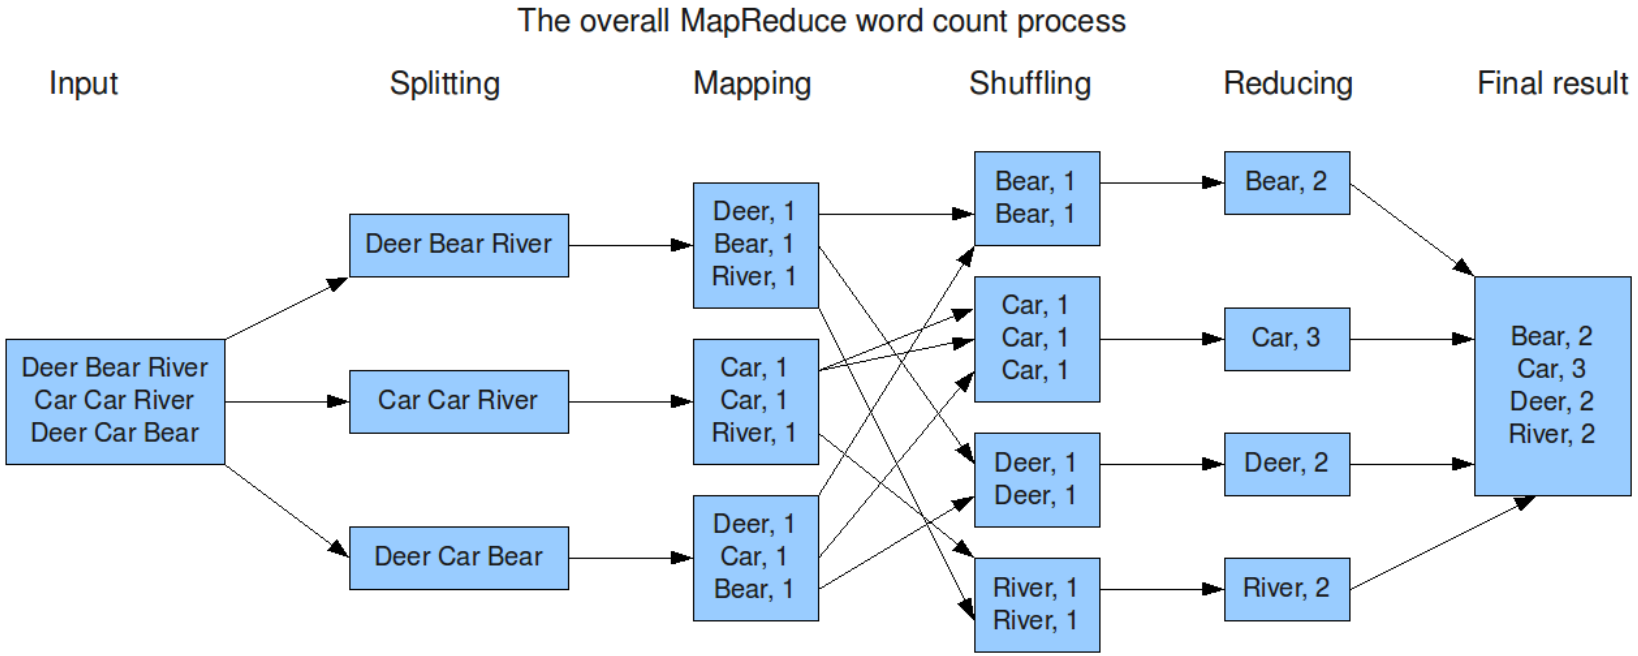
The overall MapReduce word count process
Input
Splitting
Mapping
Shuffling
Reducing
Final result
Deer Bear River
Car Car River
Deer Car Bear
Deer Bear River
Car Car River
Deer Car Bear
Deer, 1
Bear, 1
River, 1
Car, 1
Car, 1
River, 1
Deer, 1
Car, 1
Bear, 1
Bear, 1
Bear, 1
Car, 1
Car, 1
Car, 1
Deer, 1
Deer, 1
River, 1
River, 1
Bear, 2
Car, 3
Deer, 2
River, 2
Bear, 2
Car, 3
Deer, 2
River, 2

Word Count


```java
import java.io.IOException;

import java.util.Iterator;
import java.util.StringTokenizer;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reporter;
import org.apache.hadoop.mapred.TextInputFormat;
import org.apache.hadoop.mapred.TextOutputFormat;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.Reducer;

public class WordCount {

    public static class TokenizerMapper extends MapReduceBase implements
            Mapper<Object, Text, Text, IntWritable> {

        private final static IntWritable one = new IntWritable(1);

        private Text word = new Text();

        @Override
        public void map(Object key, Text value,
                OutputCollector<Text, IntWritable> output, Reporter reporter)
                throws IOException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                output.collect(word, one);
            }

        }
    }
```

```java
public static class IntSumReducer extends MapReduceBase implements
        Reducer<Text, IntWritable, Text, IntWritable> {
    private IntWritable result = new IntWritable();

    @Override
    public void reduce(Text key, Iterator<IntWritable> values,
            OutputCollector<Text, IntWritable> output, Reporter reporter)
            throws IOException {
        int sum = 0;
        while (values.hasNext()) {
            sum += values.next().get();
        }
        result.set(sum);
        output.collect(key, result);
    }
}

public static void main(String[] args) throws Exception {

    String input = "hdfs://192.168.0.110:9000/input/results.txt";
    String output = "hdfs://192.168.0.110:9000/outputs";
    JobConf conf = new JobConf(WordCount.class);
    conf.setJobName("WordCount");

    conf.setOutputKeyClass(Text.class);
    conf.setOutputValueClass(IntWritable.class);

    conf.setMapperClass(TokenizerMapper.class);
    conf.setCombinerClass(IntSumReducer.class);
    conf.setReducerClass(IntSumReducer.class);
    conf.setInputFormat(TextInputFormat.class);
    conf.setOutputFormat(TextOutputFormat.class);
    FileInputFormat.setInputPaths(conf, new Path(input));// 路径1
    FileOutputFormat.setOutputPath(conf, new Path(output));// 输出路径
    JobClient.runJob(conf);
    System.exit(0);
}
```

MapReduce API

Table 7-2. Configuration of MapReduce types in the old API

Property	JobConfsetter method	Input types		Intermediate types		Output types	
		K1	V1	K2	V2	K3	V3
Properties for configuring types:							
mapred.input.format.class	setInputFormat()	•	•				
mapred.mapoutput.key.class	setMapOutputKeyClass()			•			
mapred.mapoutput.value.class	setMapOutputValueClass()				•		
mapred.output.key.class	setOutputKeyClass()					•	
mapred.output.value.class	setOutputValueClass()						•
Properties that must be consistent with the types:							
mapred.mapper.class	setMapperClass()	•	•	•	•		
mapred.map.runner.class	setMapRunnerClass()	•	•	•	•		
mapred.combiner.class	setCombinerClass()			•	•		
mapred.partitioner.class	setPartitionerClass()			•	•		
mapred.output.key.comparator.class	setOutputKeyComparatorClass()			•			
mapred.output.value.groupfn.class	setOutputValueGroupingComparator()			•			
mapred.reducer.class	setReducerClass()			•	•	•	•
mapred.output.format.class	setOutputFormat()					•	•

MapReduce API

Table 7-1. Configuration of MapReduce types in the new API

Property	Job setter method	Input types		Intermediate types		Output types	
		K1	V1	K2	V2	K3	V3
Properties for configuring types:							
mapreduce.job.inputformat.class	setInputFormatClass()	•	•				
mapreduce.map.output.key.class	setMapOutputKeyClass()			•			
mapreduce.map.output.value.class	setMapOutputValueClass()				•		
mapreduce.job.output.key.class	setOutputKeyClass()					•	
mapreduce.job.output.value.class	setOutputValueClass()						•
Properties that must be consistent with the types:							
mapreduce.job.map.class	setMapperClass()	•	•	•	•		
mapreduce.job.combine.class	setCombinerClass()			•	•		
mapreduce.job.partitioner.class	setPartitionerClass()			•	•		
mapreduce.job.output.key.comparator.class	setSortComparatorClass()			•			
mapreduce.job.output.group.comparator.class	setGroupingComparatorClass()			•			
mapreduce.job.reduce.class	setReducerClass()			•	•	•	•
mapreduce.job.outputformat.class	setOutputFormatClass()					•	•

Use Case

KPI

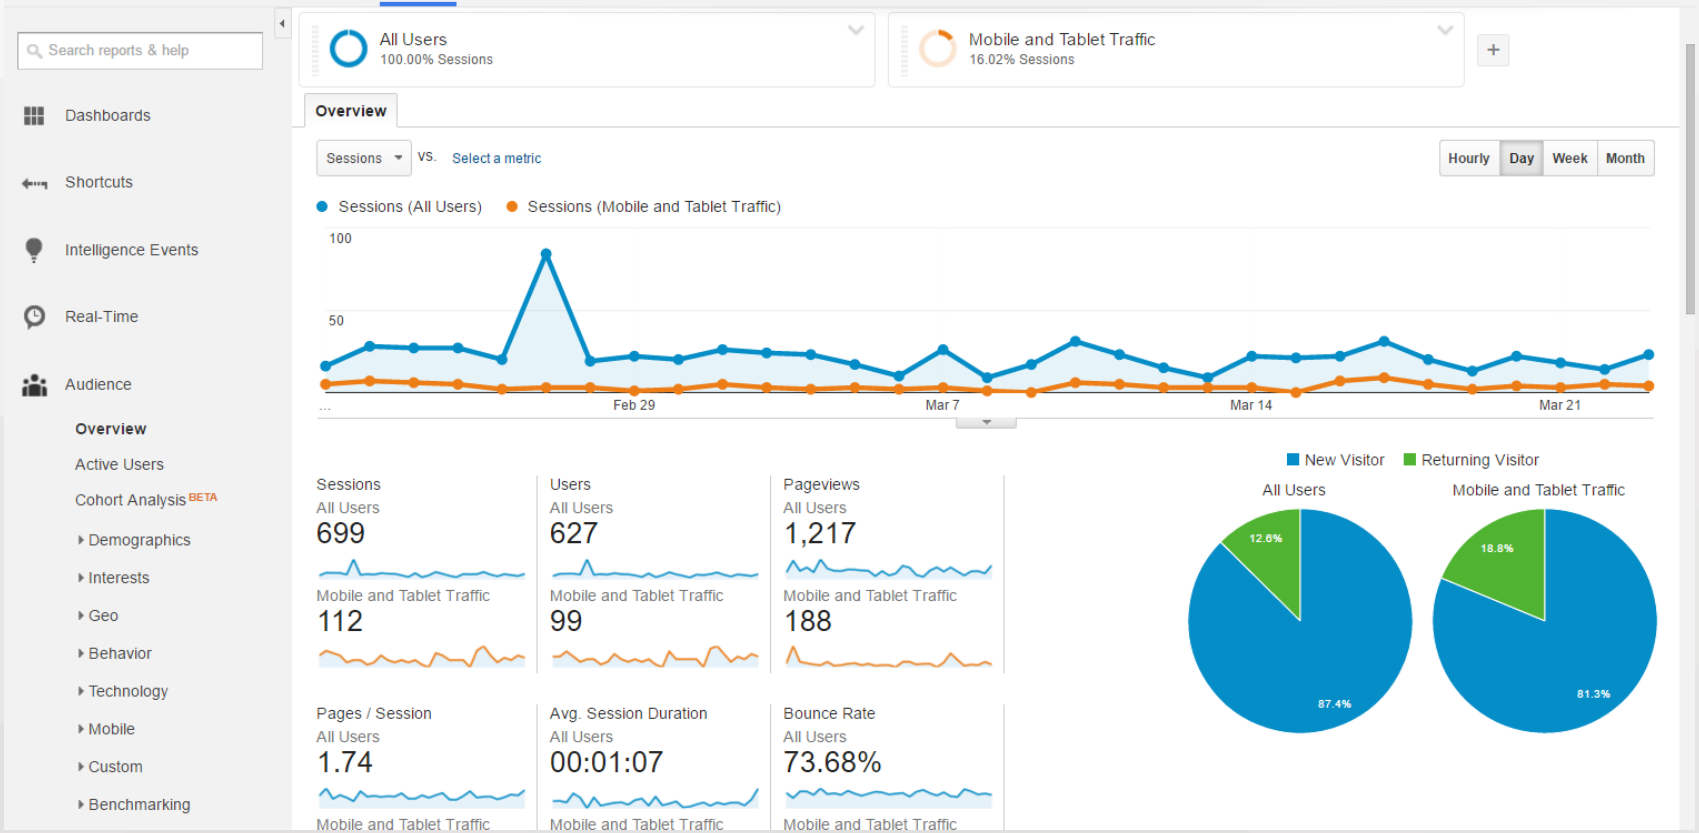
Search reports & help
Dashboards
Shortcuts
Intelligence Events
Real-Time
Audience
Overview
Active Users
Cohort Analysis BETA
Demographics
Interests
Geo
Behavior
Technology
Mobile
Custom
Benchmarking
All Users
100.00% Sessions
Mobile and Tablet Traffic
16.02% Sessions
Overview
Sessions
vs.
Select a metric
Hourly
Day
Week
Month
Sessions (All Users)
Sessions (Mobile and Tablet Traffic)
100
50
Feb 29
Mar 7
Mar 14
Mar 21
Sessions
All Users
699
Mobile and Tablet Traffic
112
Users
All Users
627
Mobile and Tablet Traffic
99
Pageviews
All Users
1,217
Mobile and Tablet Traffic
188
Pages / Session
All Users
1.74
Mobile and Tablet Traffic
Avg. Session Duration
All Users
00:01:07
Mobile and Tablet Traffic
Bounce Rate
All Users
73.68%
Mobile and Tablet Traffic
New Visitor
Returning Visitor
All Users
12.6%
87.4%
Mobile and Tablet Traffic
18.8%
81.3%

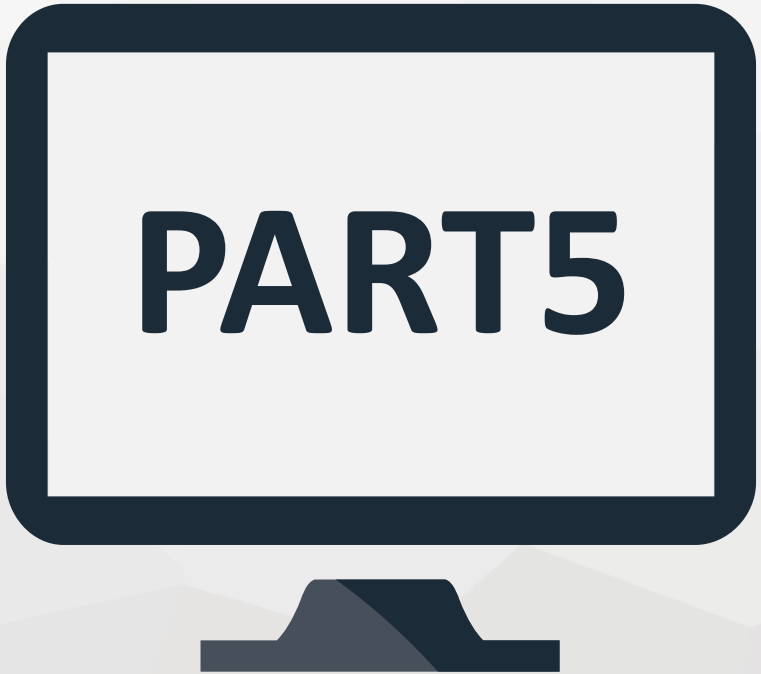

Reference Books

Books

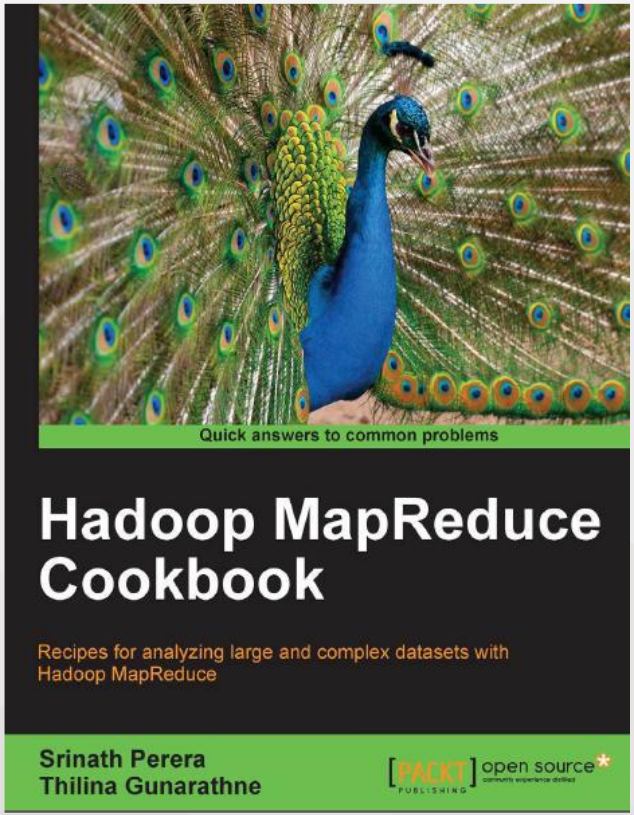
Quick answers to common problems
Hadoop MapReduce Cookbook
Recipes for analyzing large and complex datasets with Hadoop MapReduce
Srinath Perera
Thilina Gunarathne
PACKT open source

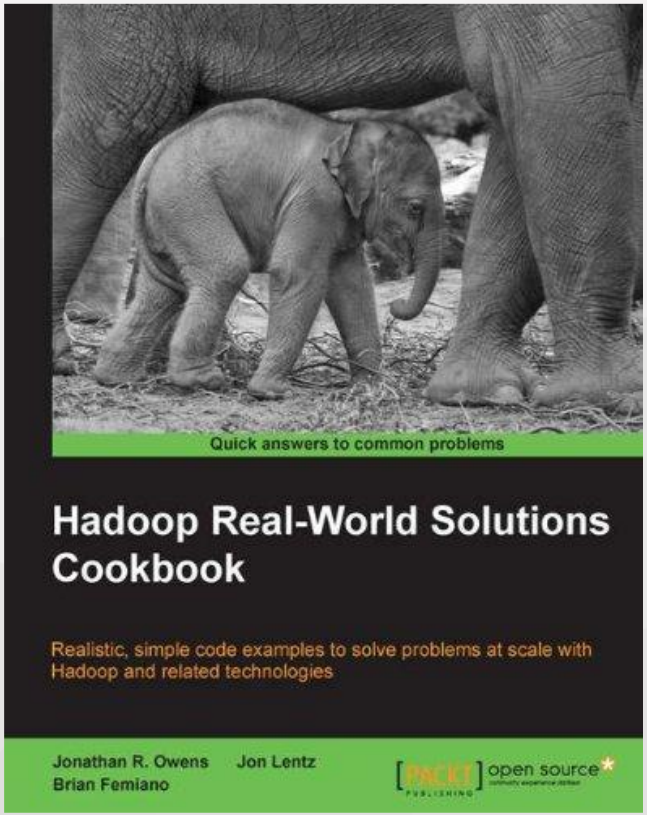
Quick answers to common problems
Hadoop Real-World Solutions Cookbook
Realistic, simple code examples to solve problems at scale with Hadoop and related technologies
Jonathan R. Owens
Jon Lentz
Brian Femiano
PACKT open source


The End