

Spark Computational Model

2017.4 XenRon

CONTENTS

Spark RDD

01

Parallel Computing


02

Computational Model

03

Spark SQL

04

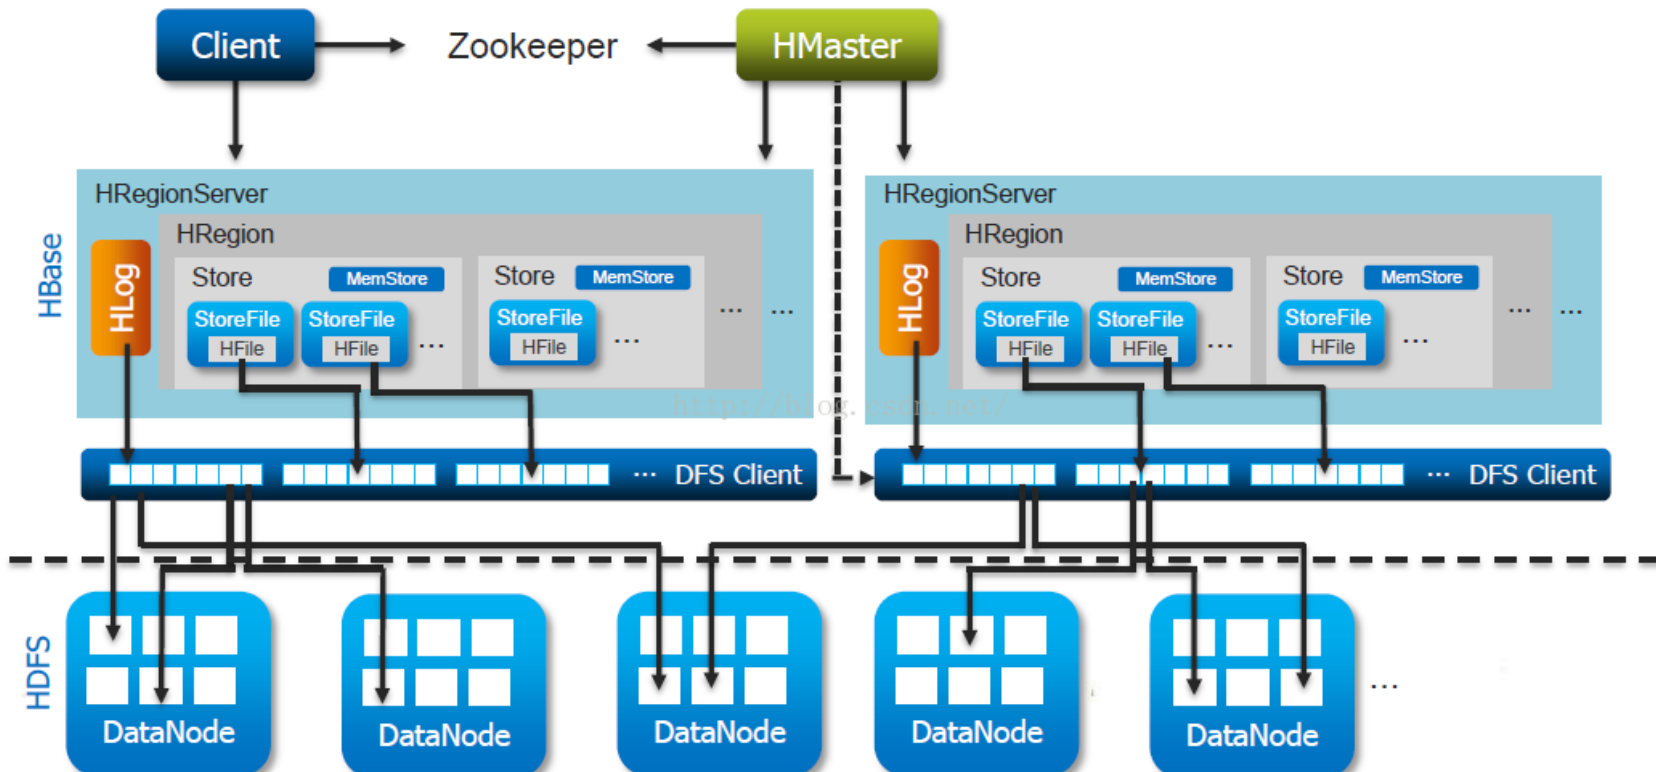





Review



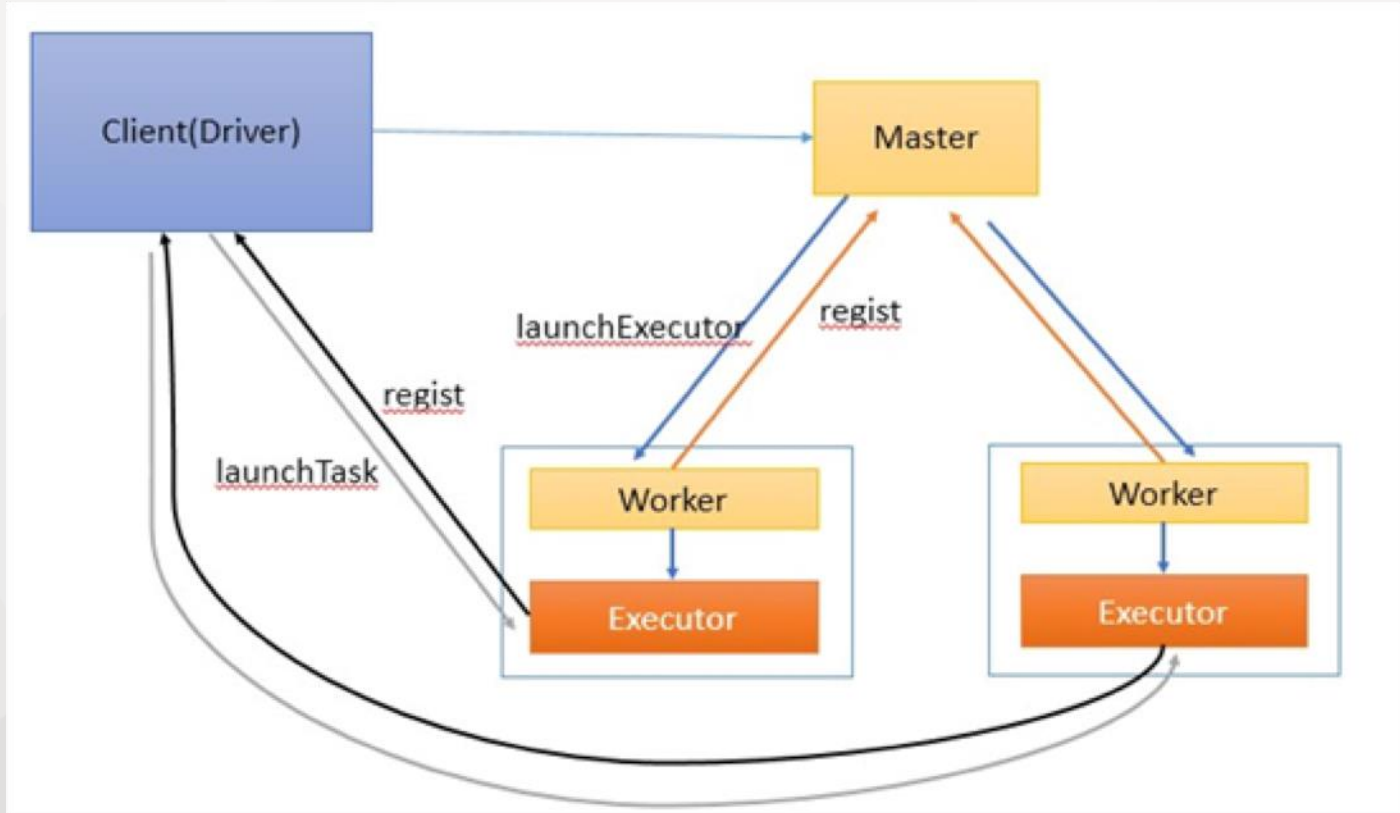
Review

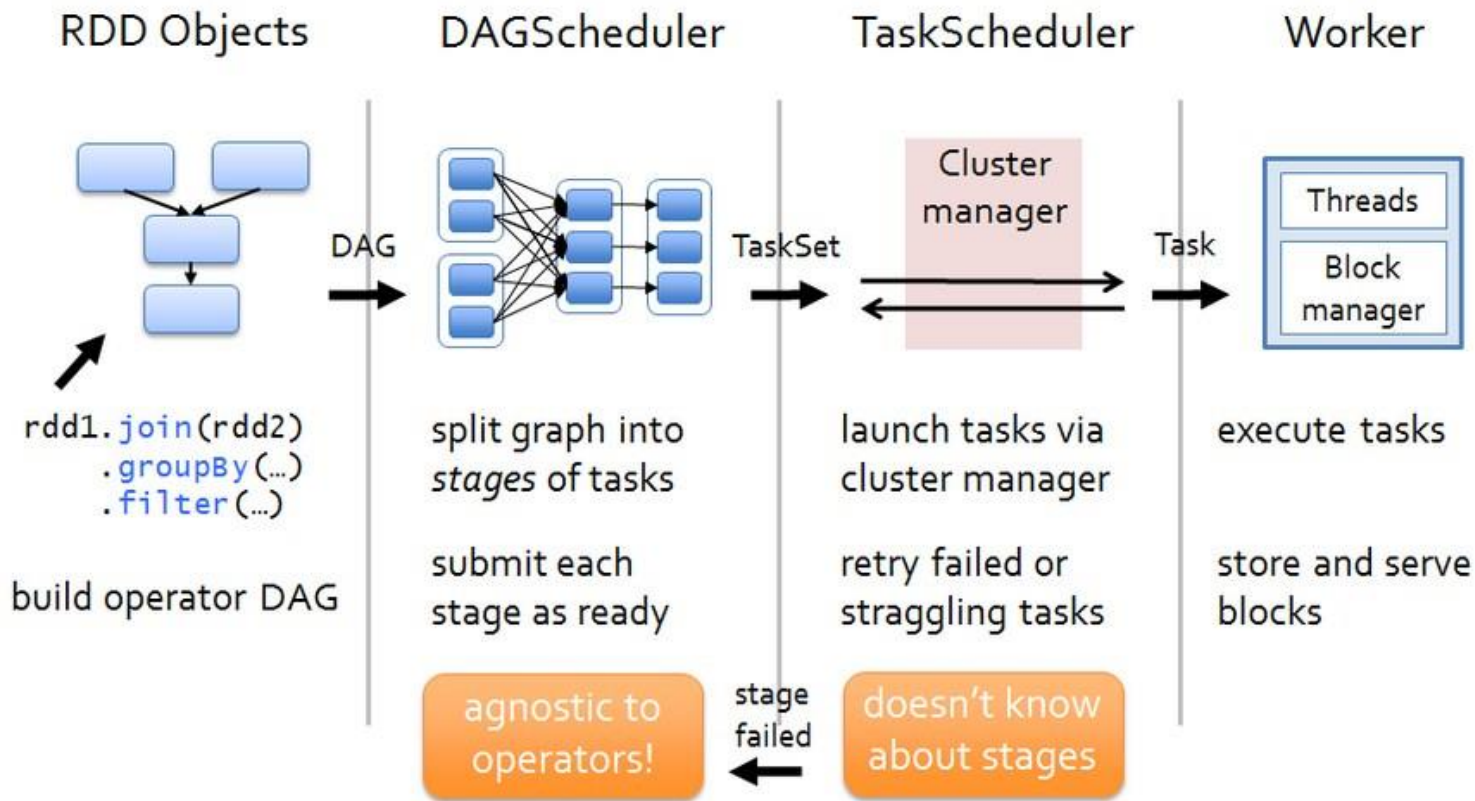


HBase high-level architecture



Spark Standalone








PART1



Spark RDD

What is an RDD?




Some RDD Characteristics


- Hold references to Partition objects
- Each Partition object references a subset of your data
- Partitions are assigned to nodes on your cluster
- Each partition/split will be in RAM (by default)

Dependency Types


“Narrow” (pipeline-able)



map, filter




union




join with inputs
co-partitioned

“Wide” (shuffle)

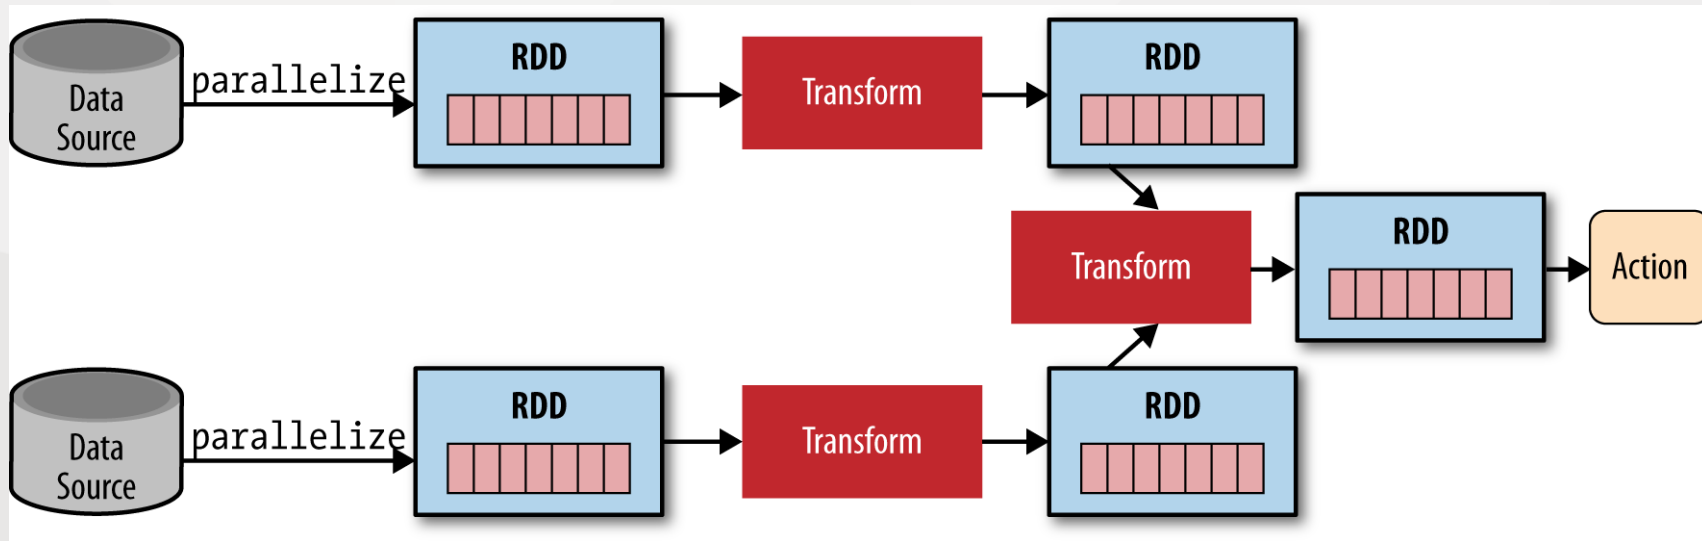


groupByKey on
non-partitioned data



join with inputs not
co-partitioned

Dependency Type



Transformations	$map(f : T \Rightarrow U) : RDD[T] \Rightarrow RDD[U]$ $filter(f : T \Rightarrow Bool) : RDD[T] \Rightarrow RDD[T]$ $flatMap(f : T \Rightarrow Seq[U]) : RDD[T] \Rightarrow RDD[U]$ $sample(fraction : Float) : RDD[T] \Rightarrow RDD[T]$ (Deterministic sampling) $groupByKey() : RDD[(K, V)] \Rightarrow RDD[(K, Seq[V])]$ $reduceByKey(f : (V, V) \Rightarrow V) : RDD[(K, V)] \Rightarrow RDD[(K, V)]$ $union() : (RDD[T], RDD[T]) \Rightarrow RDD[T]$ $join() : (RDD[(K, V)], RDD[(K, W)]) \Rightarrow RDD[(K, (V, W))]$ $cogroup() : (RDD[(K, V)], RDD[(K, W)]) \Rightarrow RDD[(K, (Seq[V], Seq[W]))]$ $crossProduct() : (RDD[T], RDD[U]) \Rightarrow RDD[(T, U)]$ $mapValues(f : V \Rightarrow W) : RDD[(K, V)] \Rightarrow RDD[(K, W)]$ (Preserves partitioning) $sort(c : Comparator[K]) : RDD[(K, V)] \Rightarrow RDD[(K, V)]$ $partitionBy(p : Partitioner[K]) : RDD[(K, V)] \Rightarrow RDD[(K, V)]$
Actions	$count() : RDD[T] \Rightarrow Long$ $collect() : RDD[T] \Rightarrow Seq[T]$ $reduce(f : (T, T) \Rightarrow T) : RDD[T] \Rightarrow T$ $lookup(k : K) : RDD[(K, V)] \Rightarrow Seq[V]$ (On hash/range partitioned RDDs) $save(path : String) : \text{Outputs RDD to a storage system, e.g., HDFS}$


Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(_.startsWith("ERROR"))
messages = errors.map(_.split('\t')(2))
cachedMsgs = messages.cache()
```

```
cachedMsgs.filter(_.contains("foo")).count
cachedMsgs.filter(_.contains("bar")).count
. . .
```

Result: scaled to 1 TB data in 5-7 sec
(vs 170 sec for on-disk data)







PART2



Parallel Computing

Parallel Computing

15



Parallel Computing



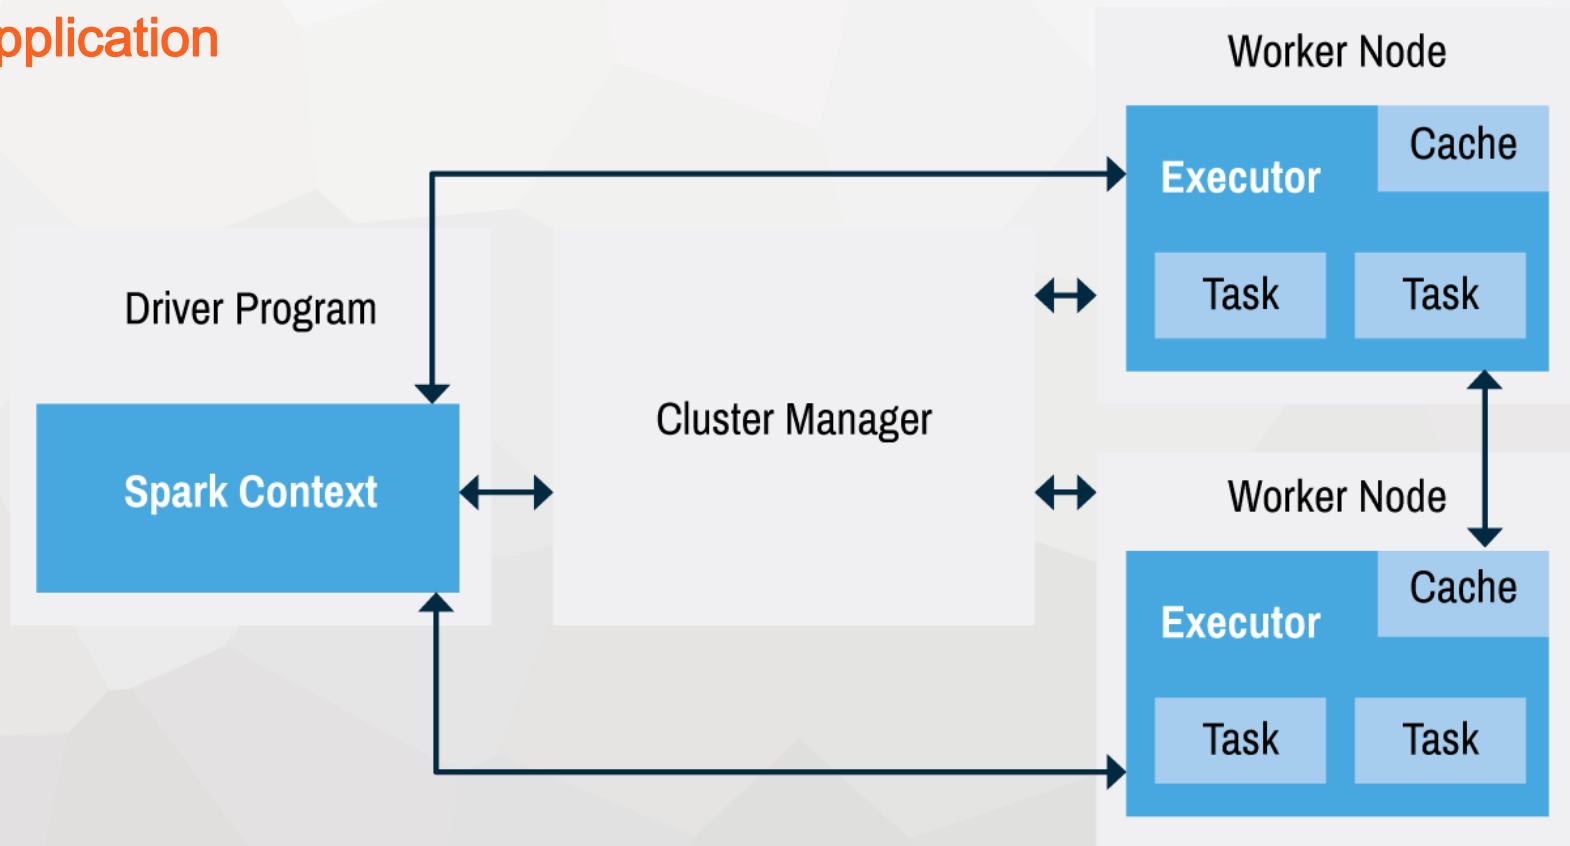


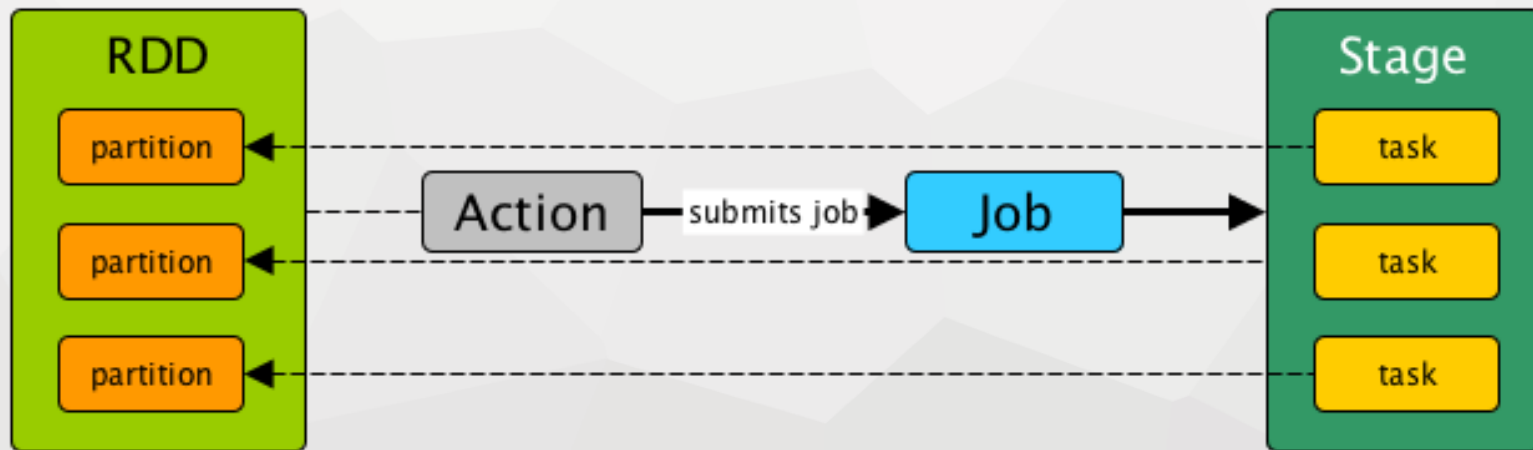
PART3

Computational Model

Driver Executor Application

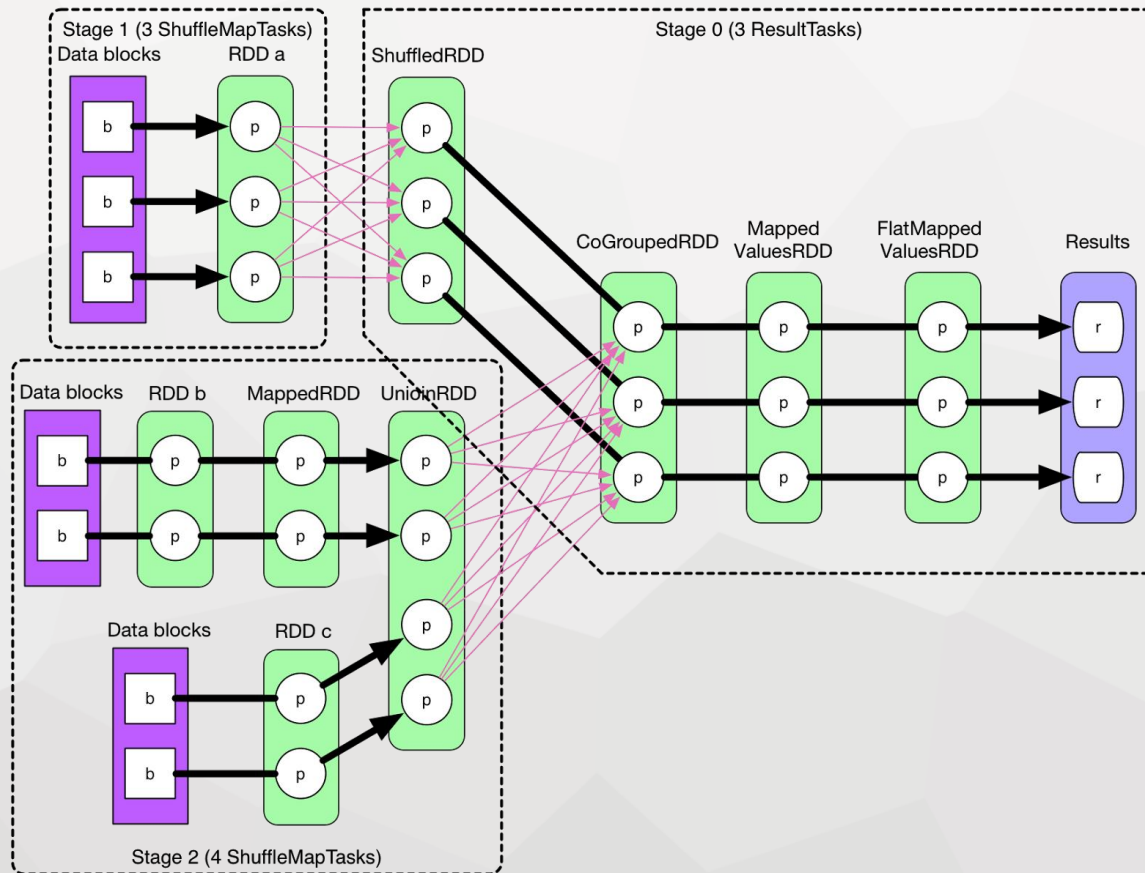
18

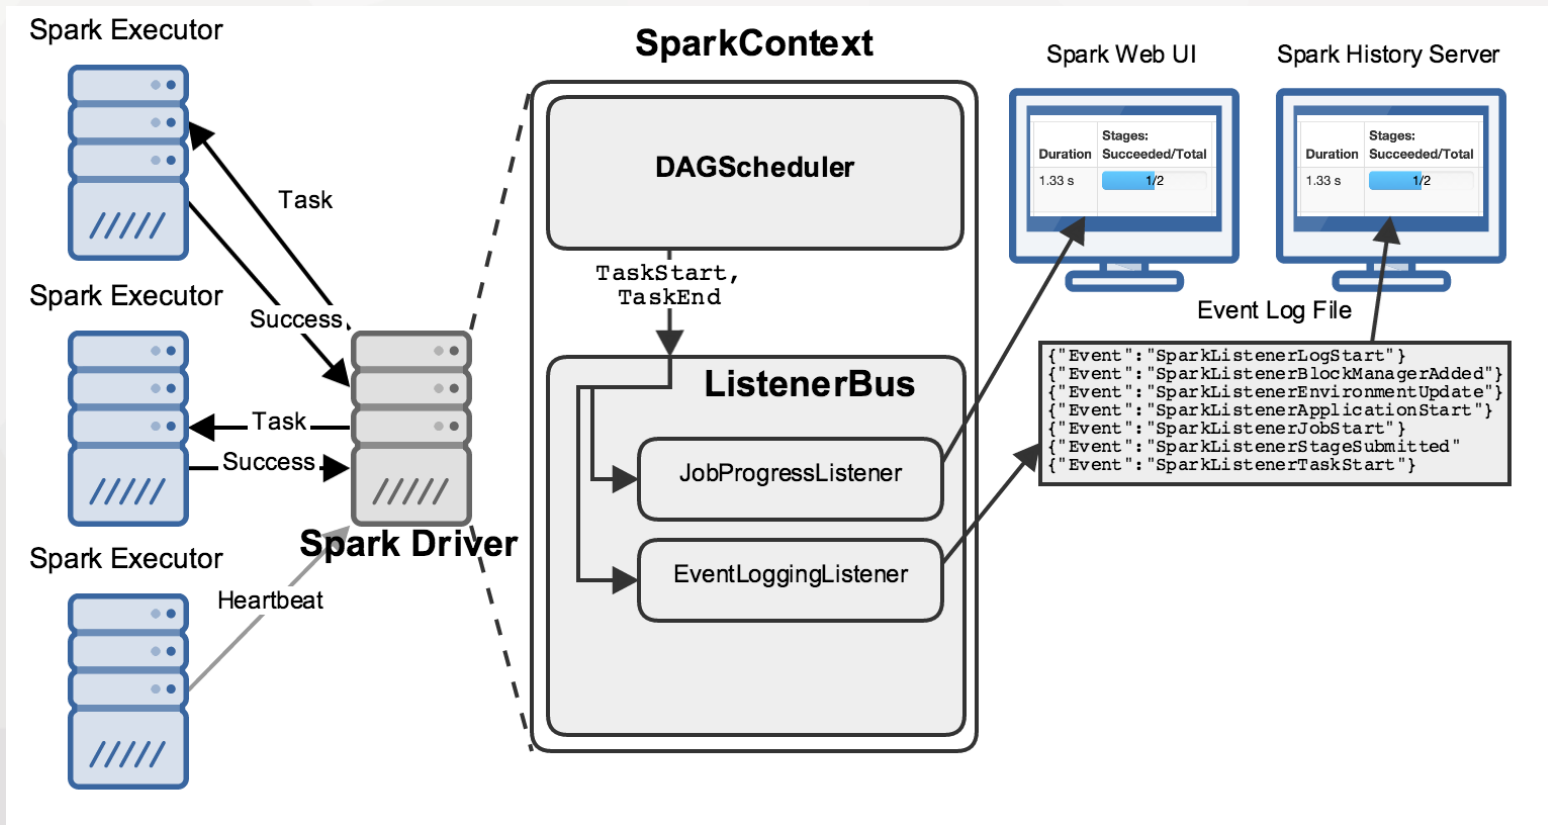




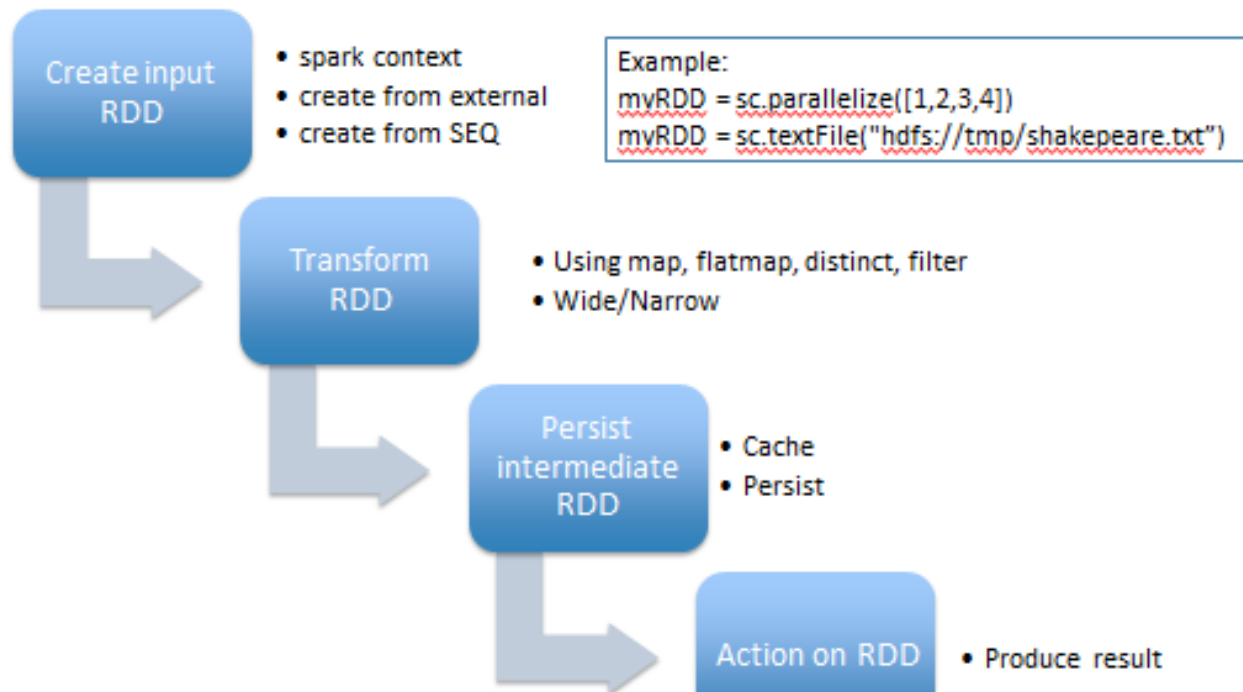
Stage

ComplexJob
including map(), partitionBy(), union(), and join()



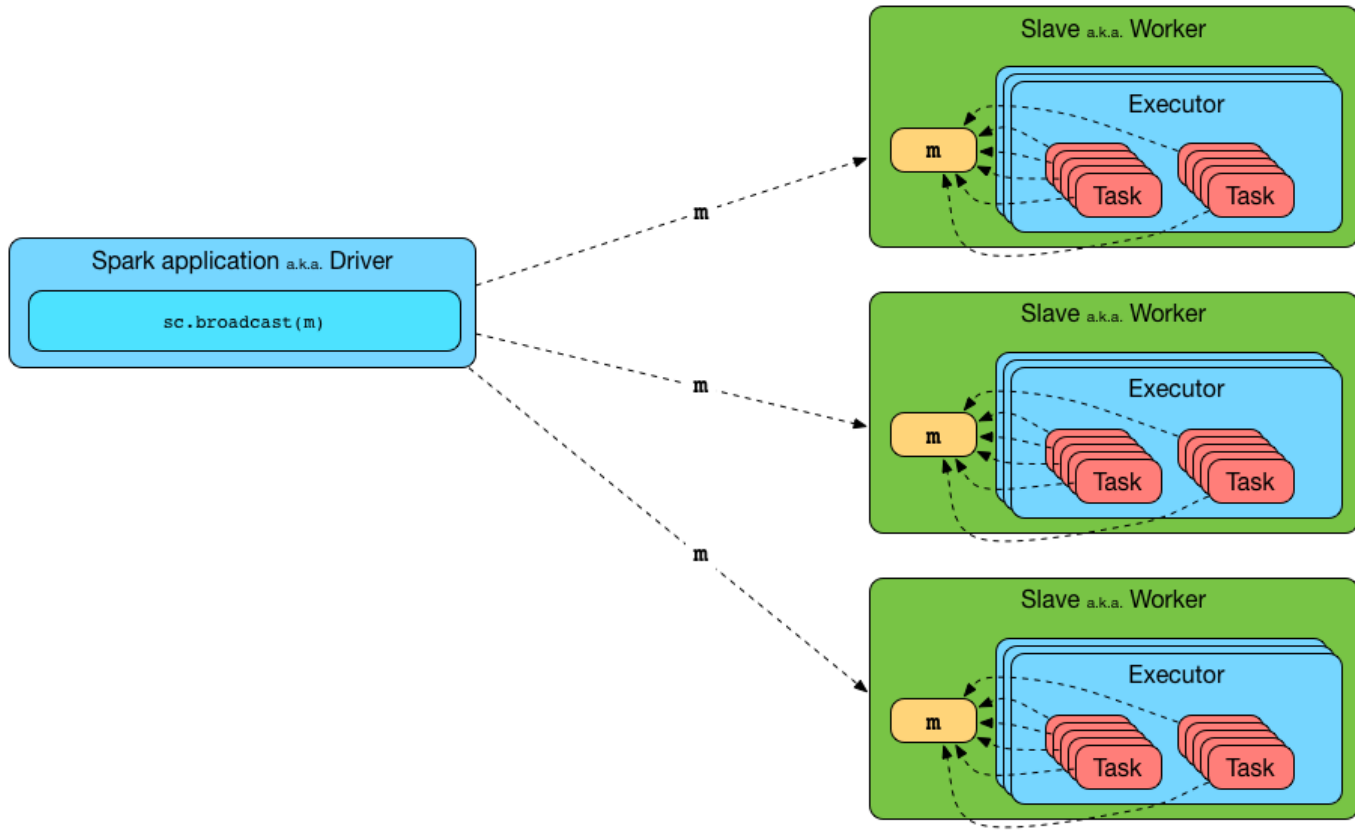


Spark Program Flow by RDD





Broad Cast



ACCUMULATORS

- variables that are only “added” to through an associative operation (add())
- only the driver program can read the accumulator's value

```
Accumulator<Integer> accum = sc.accumulator(0);  
  
sc.parallelize(Arrays.asList(1, 2, 3, 4)).foreach(x ->  
    accum.add(x));  
  
accum.value();  
// returns 10
```


Word Count

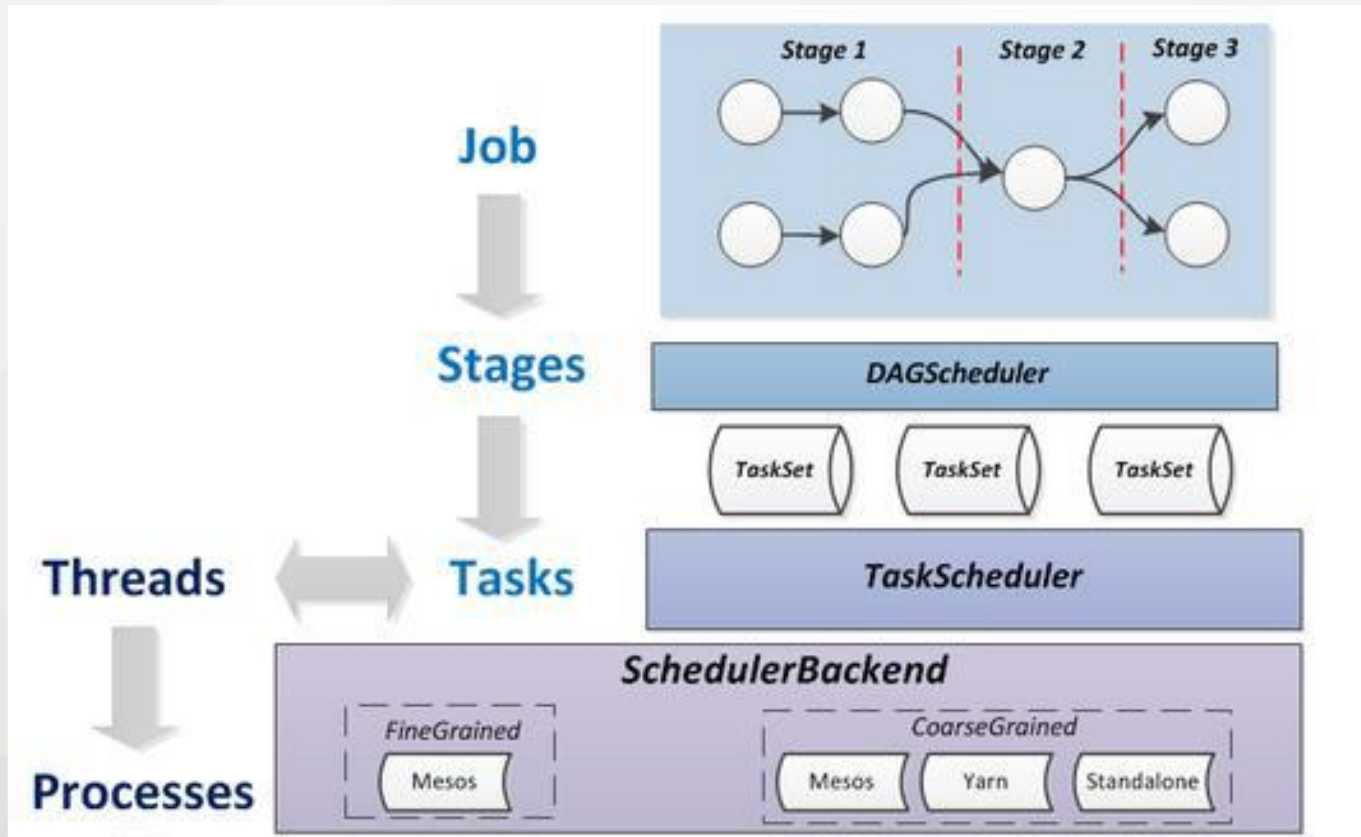
```
import sys
from pyspark import SparkContext
from pyspark.mllib.regression import LabeledPoint, LinearRegressionWithSGD
from numpy import array

# Load and parse the data
def parsePoint(line):
    values = [float(x) for x in line.replace(',', ' ').split(' ')]
    return LabeledPoint(values[0], values[1:])

sc = SparkContext(appName="LinearRegressionPredict")
data = sc.textFile("data/mllib/ridge-data/lpsa.data")
parsedData = data.map(parsePoint)

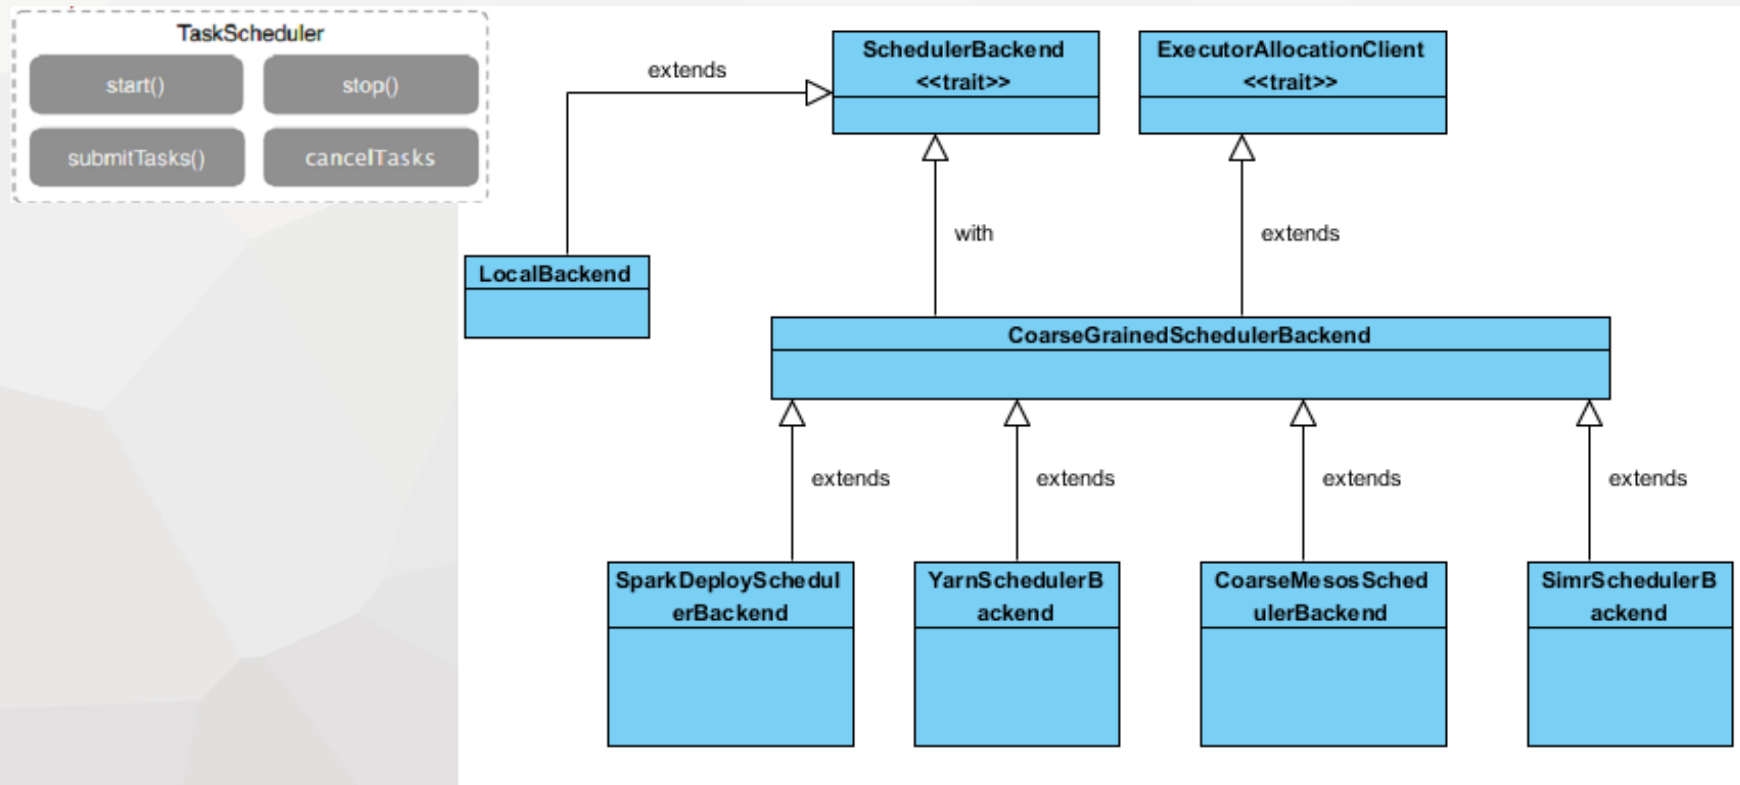
# Build the model
model = LinearRegressionWithSGD.train(parsedData)

# Evaluate the model on training data
valuesAndPreds = parsedData.map(lambda p: (p.label, model.predict(p.features)))
MSE = valuesAndPreds.map(lambda (v, p): (v - p)**2).reduce(lambda x, y: x + y) /
valuesAndPreds.count()
print("Mean Squared Error = " + str(MSE))
sc.stop()
```



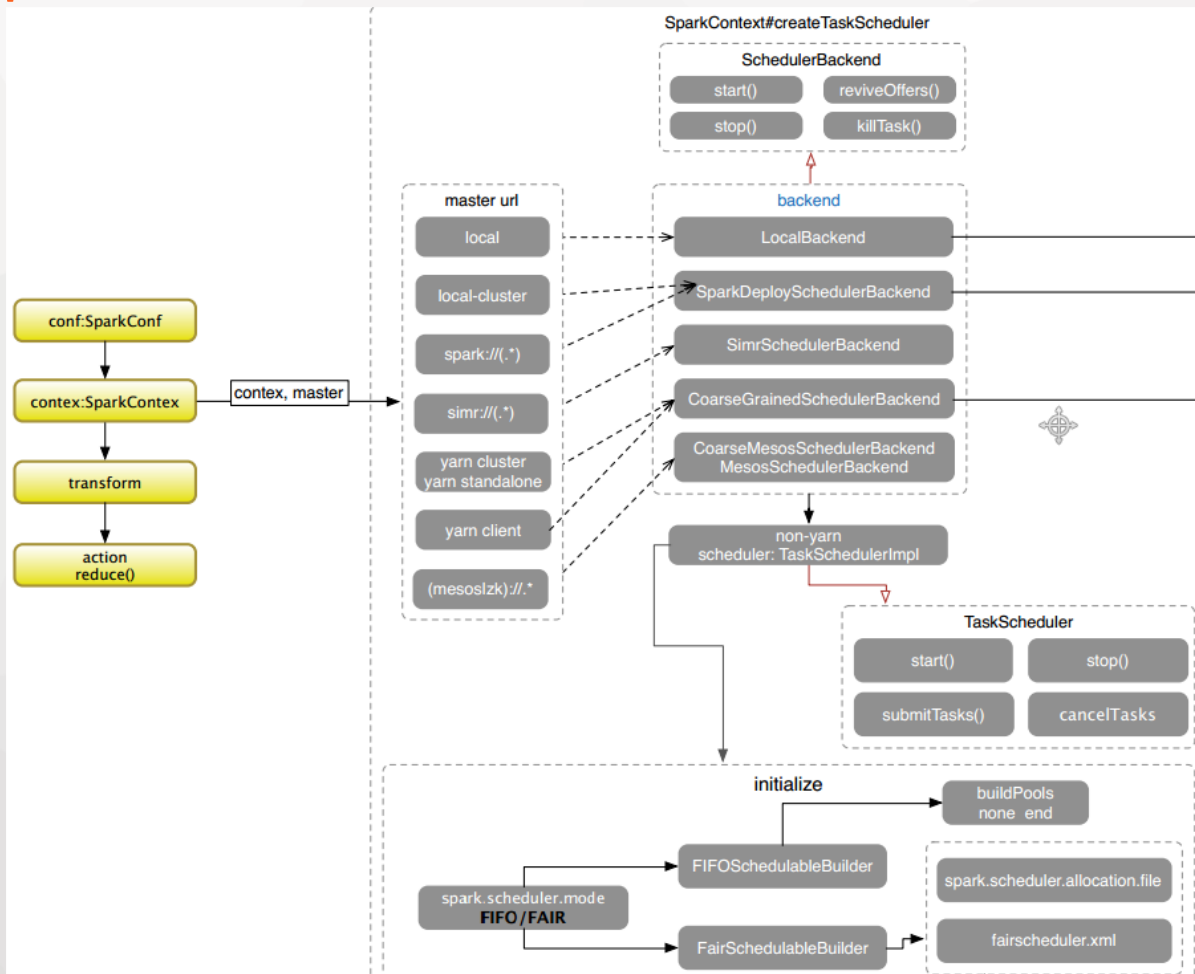
TaskScheduler

27



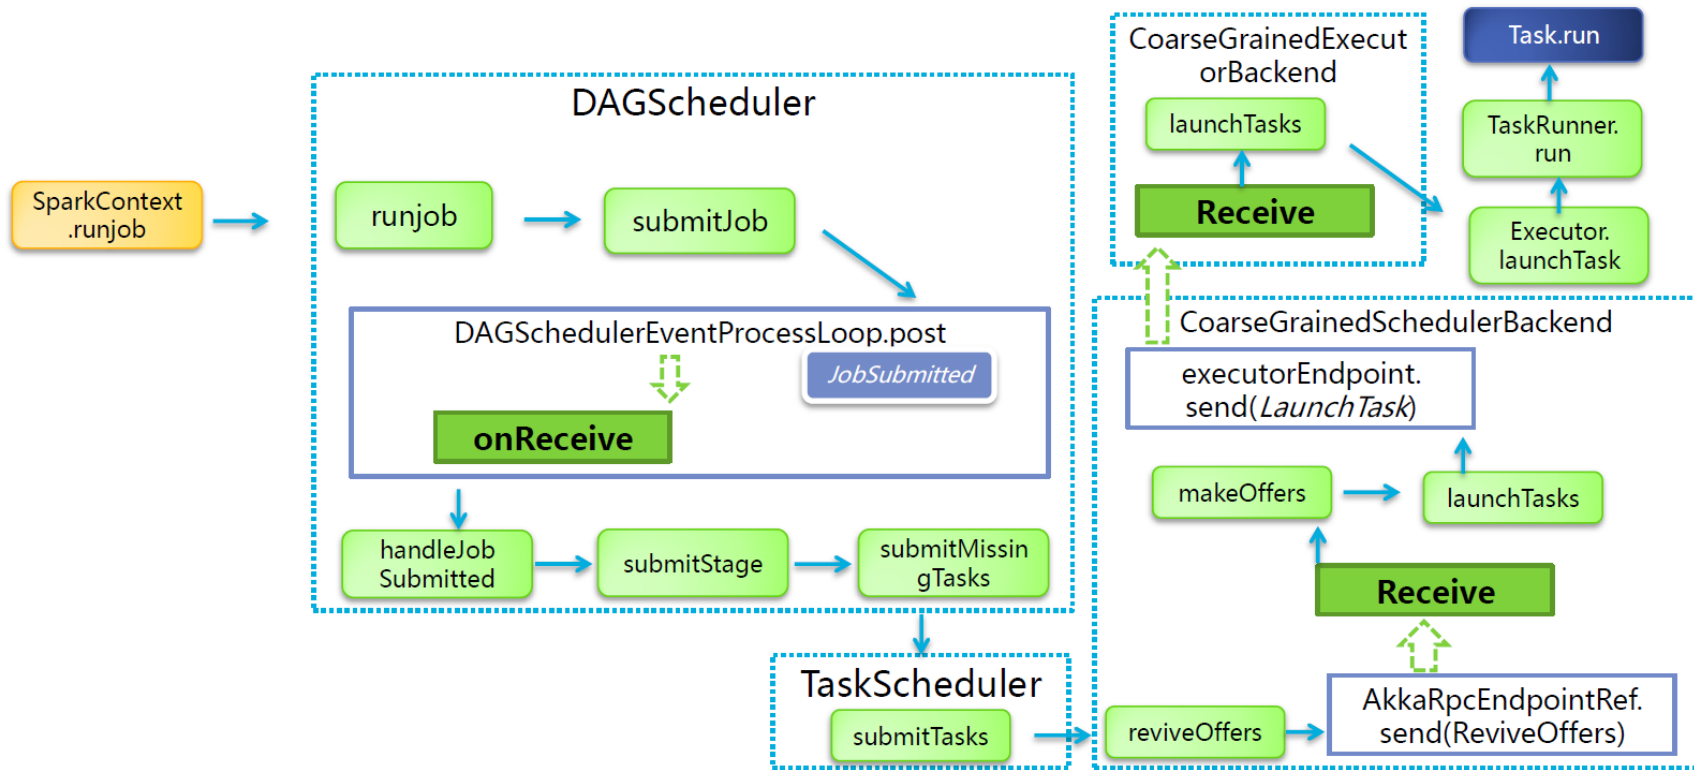
TaskScheduler

28



TaskScheduler

29





PART4




Spark SQL

Hive



- Data warehousing package built on top of Hadoop
- Bringing structure to unstructured data
- Query petabytes of data with HiveQL
- Schema on read



Hadoop MapReduce Vs Pig Vs Hive

Hadoop MapReduce	Pig	Hive
Compiled Language	Scripting Language	SQL like query Language
Lower Level of Abstraction	Higher Level of Abstraction	Higher Level of Abstraction
More lines of Code	Comparatively less lines of Code than MapReduce	Comparatively less lines of Code than MapReduce and Apache Pig
More Development Effort is involved	Development Effort is less Code Efficiency is relatively less	Development Effort is less Code Efficiency is relatively less
Code Efficiency is high when compared to Pig and Hive	Code Efficiency is relatively less	Code Efficiency is relatively less

The Right SQL Engine for the Use Case



BI and SQL
Analytics

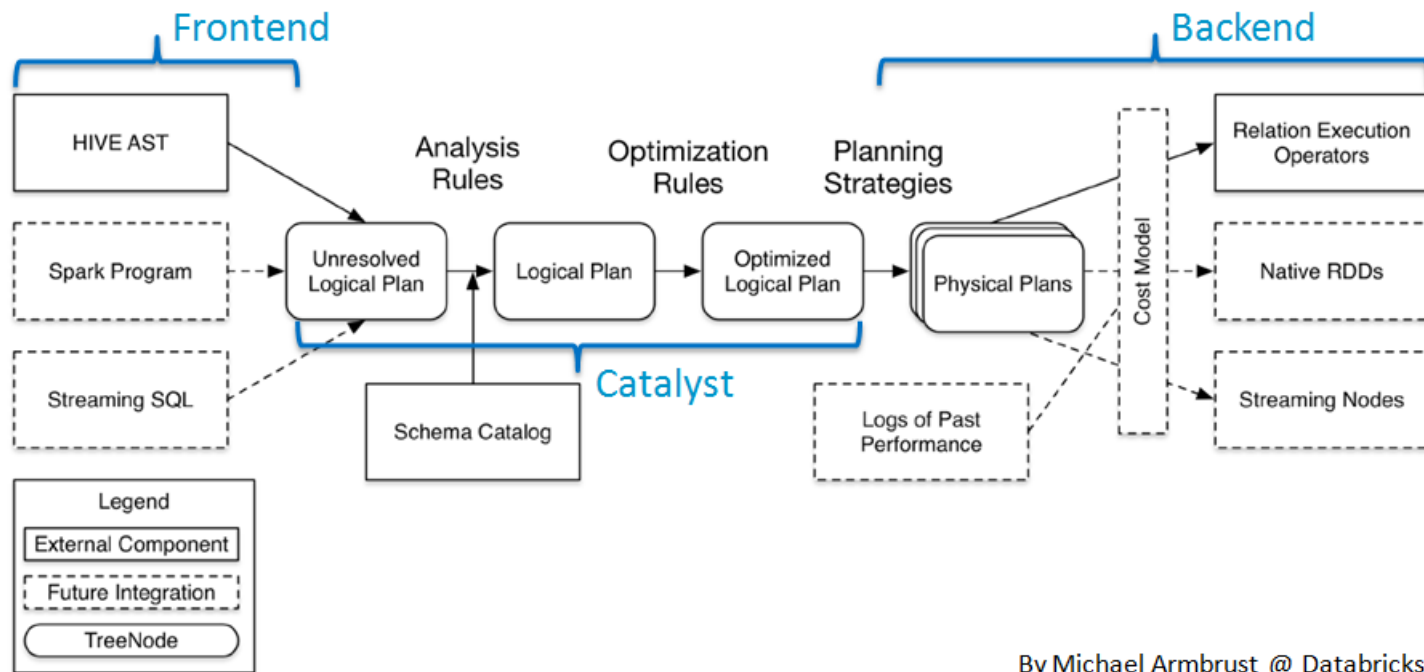


Batch
Processing



Spark
Developers

Spark SQL Architecture



By Michael [Armbrust](#) @ [Databricks](#)



The End