

HBase shell commands

As told in HBase introduction, HBase provides Extensible jruby-based (JIRB) shell as a feature to execute some commands(each command represents one functionality).

HBase shell commands are mainly categorized into 6 parts

1) General HBase shell commands

status

Show cluster status. Can be 'summary', 'simple', or 'detailed'. The default is 'summary'.

```
hbase> status
```

```
hbase> status 'simple'
```

```
hbase> status 'summary'
```

```
hbase> status 'detailed'
```

version

Output this HBase versionUsage:

```
hbase> version
```

whoami

Show the current hbase user.Usage:

```
hbase> whoami
```

2) Tables Management commands

alter

Alter column family schema; pass table name and a dictionary specifying new column family schema. Dictionaries are described on the main help command output. Dictionary must include name of column family to alter.For example, to change or add the 'f1' column family in table 't1' from

current value to keep a maximum of 5 cell VERSIONS, do:

```
hbase> alter 't1', NAME => 'f1', VERSIONS => 5
```

You can operate on several column families:

```
hbase> alter 't1', 'f1', {NAME => 'f2', IN_MEMORY => true}, {NAME => 'f3', VERSIONS => 5}
```

To delete the 'f1' column family in table 't1', use one of:
hbase> alter 't1', NAME => 'f1',
METHOD => 'delete'

hbase> alter 't1', 'delete' => 'f1'

You can also change table-scope attributes like MAX_FILESIZE, READONLY, MEMSTORE_FLUSH_SIZE, DEFERRED_LOG_FLUSH, etc. These can be put at the end; for example, to change the max size of a region to 128MB, do:

hbase> alter 't1', MAX_FILESIZE => '134217728'

You can add a table coprocessor by setting a table coprocessor attribute:

hbase> alter 't1',
'coprocessor'=>'hdfs:///foo.jar|com.foo.FooRegionObserver|1001|arg1=1,arg2=2'

Since you can have multiple coprocessors configured for a table, a sequence number will be automatically appended to the attribute name to uniquely identify it.

The coprocessor attribute must match the pattern below in order for the framework to understand how to load the coprocessor classes:

[coprocessor jar file location] | class name | [priority] | [arguments]

You can also set configuration settings specific to this table or column family:

hbase> alter 't1', CONFIGURATION =>
{ 'hbase.hregion.scan.loadColumnFamiliesOnDemand' => 'true' }
hbase> alter 't1', { NAME => 'f2', CONFIGURATION => { 'hbase.hstore.blockingStoreFiles' => '10' } }

You can also remove a table-scope attribute:

hbase> alter 't1', METHOD => 'table_att_unset', NAME => 'MAX_FILESIZE'
hbase> alter 't1', METHOD => 'table_att_unset', NAME => 'coprocessor\$1'

There could be more than one alteration in one command:

hbase> alter 't1', { NAME => 'f1', VERSIONS => 3 },
{ MAX_FILESIZE => '134217728' }, { METHOD => 'delete', NAME => 'f2' },
OWNER => 'johndoe', METADATA => { 'mykey' => 'myvalue' }

create

Create table; pass table name, a dictionary of specifications per column family, and optionally a dictionary of table configuration.

hbase> create 't1', { NAME => 'f1', VERSIONS => 5 }
hbase> create 't1', { NAME => 'f1' }, { NAME => 'f2' }, { NAME => 'f3' }

hbase> # The above in shorthand would be the following:

hbase> create 't1', 'f1', 'f2', 'f3'

hbase> create 't1', {NAME => 'f1', VERSIONS => 1, TTL => 2592000, BLOCKCACHE => true}

hbase> create 't1', {NAME => 'f1', CONFIGURATION => {'hbase.hstore.blockingStoreFiles' => '10'}}

Table configuration options can be put at the end.

describe

Describe the named table.

hbase> describe 't1'

disable

Start disable of named table

hbase> disable 't1'

disable_all

Disable all of tables matching the given regex

hbase> disable_all 't.*'

is_disabled

verifies Is named table disabled

hbase> is_disabled 't1'

drop

Drop the named table. Table must first be disabled

hbase> drop 't1'

drop_all

Drop all of the tables matching the given regex

hbase> drop_all 't.*'

enable

Start enable of named table

hbase> enable 't1'

enable_all

Enable all of the tables matching the given regex

```
hbase> enable_all 't.*'
```

is_enabled

verifies if named table enabled

```
hbase> is_enabled 't1'
```

exists

Does the named table exist

```
hbase> exists 't1'
```

list

List all tables in hbase. Optional regular expression parameter could be used to filter the output

```
hbase> list
```

```
hbase> list 'abc.*'
```

show_filters

Show all the filters in hbase.

```
hbase> show_filters
```

alter_status

Get the status of the alter command. Indicates the number of regions of the table that have received the updated schema Pass table name.

```
hbase> alter_status 't1'
```

alter_async

Alter column family schema, does not wait for all regions to receive the schema changes. Pass table name and a dictionary specifying new column family schema. Dictionaries are described on the main help command output.

Dictionary must include name of column family to alter.

To change or add the 'f1' column family in table 't1' from defaults

to instead keep a maximum of 5 cell VERSIONS, do: `hbase> alter_async 't1', NAME => 'f1', VERSIONS => 5` To delete the 'f1' column family in table 't1', do:

```
hbase> alter_async 't1', NAME => 'f1', METHOD => 'delete' or a shorter version: hbase> alter_async 't1', 'delete' => 'f1'
```

You can also change table-scope attributes like MAX_FILESIZE

MEMSTORE_FLUSH_SIZE, READONLY, and DEFERRED_LOG_FLUSH.

For example, to change the max size of a family to 128MB, do:

```
hbase> alter 't1', METHOD => 'table_att', MAX_FILESIZE => '134217728'
```

There could be more than one alteration in one command:

```
hbase> alter 't1', {NAME => 'f1'}, {NAME => 'f2', METHOD => 'delete'}
```

To check if all the regions have been updated, use `alter_status`

3) Data Manipulation commands

count

Count the number of rows in a table. Return value is the number of rows.

This operation may take a LONG time (Run '\$HADOOP_HOME/bin/hadoop jar hbase.jar rowcount' to run a counting mapreduce job). Current count is shown every 1000 rows by default. Count interval may be optionally specified. Scan caching is enabled on count scans by default. Default cache size is 10 rows.

If your rows are small in size, you may want to increase this

parameter. Examples: `hbase> count 't1'`

```
hbase> count 't1', INTERVAL => 100000
```

```
hbase> count 't1', CACHE => 1000
```

```
hbase> count 't1', INTERVAL => 10, CACHE => 1000
```

The same commands also can be run on a table reference. Suppose you had a reference `t` to table 't1', the corresponding commands would be: `hbase> t.count`

```
hbase> t.count INTERVAL => 100000
```

```
hbase> t.count CACHE => 1000
```

```
hbase> t.count INTERVAL => 10, CACHE => 1000
```

delete

Put a delete cell value at specified table/row/column and optionally timestamp coordinates. Deletes must match the deleted cell's coordinates exactly. When scanning, a delete cell suppresses older versions. To delete a cell from 't1' at row 'r1' under column 'c1' marked with the time 'ts1', do:

```
hbase> delete 't1', 'r1', 'c1', ts1
```

The same command can also be run on a table reference. Suppose you had a reference `t` to table 't1', the corresponding command would be: `hbase> t.delete 'r1', 'c1', ts1`

deleteall

Delete all cells in a given row; pass a table name, row, and optionally a column and timestamp. Examples:

```
hbase> deleteall 't1', 'r1'
```

```
hbase> deleteall 't1', 'r1', 'c1', ts1
```

The same commands also can be run on a table reference. Suppose you had a reference `t` to table 't1', the corresponding command would be:

```
hbase> t.deleteall 'r1', 'c1'
```

```
hbase> t.deleteall 'r1', 'c1', ts1
```

get

Get row or cell contents; pass table name, row, and optionally a dictionary of column(s), timestamp, timerange and versions. Examples:

```
hbase> get 't1', 'r1'
```

```
hbase> get 't1', 'r1', {TIMERANGE => [ts1, ts2]}
```

```
hbase> get 't1', 'r1', {COLUMN => 'c1'}
```

```
hbase> get 't1', 'r1', {COLUMN => ['c1', 'c2', 'c3']}
```

```
hbase> get 't1', 'r1', {COLUMN => 'c1', TIMESTAMP => ts1}
```

```
hbase> get 't1', 'r1', {COLUMN => 'c1', TIMERANGE => [ts1, ts2], VERSIONS => 4}
```

```
hbase> get 't1', 'r1', {COLUMN => 'c1', TIMESTAMP => ts1, VERSIONS => 4}
```

```
hbase> get 't1', 'r1', {FILTER => "ValueFilter(=, 'binary:abc)'"}
```

```
hbase> get 't1', 'r1', 'c1'
```

```
hbase> get 't1', 'r1', 'c1', 'c2'
```

```
hbase> get 't1', 'r1', ['c1', 'c2']
```

Besides the default 'toStringBinary' format, 'get' also supports custom formatting by column. A user can define a FORMATTER by adding it to the column name in the get specification. The FORMATTER can be stipulated: 1. either as a `org.apache.hadoop.hbase.util.Bytes` method name (e.g. `toInt`, `toString`)

1. or as a custom class followed by method name: e.g.

'c(MyFormatterClass).format'. Example formatting `cf:qualifier1` and `cf:qualifier2` both as Integers:

```
hbase> get 't1', 'r1' {COLUMN => ['cf:qualifier1:toInt',  
  'cf:qualifier2:c(org.apache.hadoop.hbase.util.Bytes).toInt'] }
```

Note that you can specify a FORMATTER by column only (`cf:qualifier`). You cannot specify a FORMATTER for all columns of a column family. The same commands also can be run on a reference to a table (obtained via `get_table` or `create_table`). Suppose you had a reference `t` to table 't1', the corresponding commands would be:

```
hbase> t.get 'r1'
```

```
hbase> t.get 'r1', {TIMERANGE => [ts1, ts2]}
hbase> t.get 'r1', {COLUMN => 'c1'}
hbase> t.get 'r1', {COLUMN => ['c1', 'c2', 'c3']}
hbase> t.get 'r1', {COLUMN => 'c1', TIMESTAMP => ts1}
hbase> t.get 'r1', {COLUMN => 'c1', TIMERANGE => [ts1, ts2], VERSIONS => 4}
hbase> t.get 'r1', {COLUMN => 'c1', TIMESTAMP => ts1, VERSIONS => 4}
hbase> t.get 'r1', {FILTER => "ValueFilter(=, 'binary:abc')"}
hbase> t.get 'r1', 'c1'
hbase> t.get 'r1', 'c1', 'c2'
hbase> t.get 'r1', ['c1', 'c2']
```

get_counter

Return a counter cell value at specified table/row/column coordinates.

A cell should be managed with atomic increment function on HBase and the data should be binary encoded. Example:

```
hbase> get_counter 't1', 'r1', 'c1'
```

The same commands also can be run on a table reference. Suppose you had a reference `t` to table 't1', the corresponding command would be:

```
hbase> t.get_counter 'r1', 'c1'
```

incr

Increments a cell 'value' at specified table/row/column coordinates.

To increment a cell value in table 't1' at row 'r1' under column 'c1' by 1 (can be omitted) or 10 do:

```
hbase> incr 't1', 'r1', 'c1'
hbase> incr 't1', 'r1', 'c1', 1
hbase> incr 't1', 'r1', 'c1', 10
```

The same commands also can be run on a table reference. Suppose you had a reference `t` to table 't1', the corresponding command would be:

```
hbase> t.incr 'r1', 'c1'
hbase> t.incr 'r1', 'c1', 1
hbase> t.incr 'r1', 'c1', 10
```

put

Put a cell 'value' at specified table/row/column and optionally timestamp coordinates. To put a cell value into table 't1' at row 'r1' under column 'c1' marked with the time 'ts1', do:

```
hbase> put 't1', 'r1', 'c1', 'value', ts1
```

The same commands also can be run on a table reference. Suppose you had a reference `t` to table 't1', the corresponding command would be:

```
hbase> t.put 'r1', 'c1', 'value', ts1
```

scan

Scan a table; pass table name and optionally a dictionary of scanner

specifications. Scanner specifications may include one or more of:

TIMERANGE, FILTER, LIMIT, STARTROW, STOPROW, TIMESTAMP, MAXLENGTH, or COLUMNS. CACHE. If no columns are specified, all columns will be scanned.

To scan all members of a column family, leave the qualifier empty as in

'col_family:'. The filter can be specified in two ways:

1. Using a filterString – more information on this is available in the Filter Language document attached to the HBASE-4176 JIRA
2. Using the entire package name of the filter. Some examples:

```
hbase> scan '.META.'
```

```
hbase> scan '.META.', {COLUMNS => 'info:regioninfo'}
```

```
hbase> scan 't1', {COLUMNS => ['c1', 'c2'], LIMIT => 10, STARTROW => 'xyz'}
```

```
hbase> scan 't1', {COLUMNS => 'c1', TIMERANGE => [1303668804, 1303668904]}
```

```
hbase> scan 't1', {FILTER => "(PrefixFilter ('row2') AND  
(QualifierFilter (>=, 'binary:xyz')) AND (TimestampsFilter ( 123, 456)))"}
```

```
hbase> scan 't1', {FILTER =>  
org.apache.hadoop.hbase.filter.ColumnPaginationFilter.new(1, 0)}
```

For experts, there is an additional option — `CACHE_BLOCKS` — which switches block caching for the scanner on (true) or off (false). By default it is enabled. Examples:

```
hbase> scan 't1', {COLUMNS => ['c1', 'c2'],  
CACHE_BLOCKS => false}
```

Also for experts, there is an advanced option — `RAW` — which instructs the scanner to return all cells (including delete markers and uncollected deleted cells). This option cannot be combined with requesting specific COLUMNS.

Disabled by default. Example:

```
hbase> scan 't1', {RAW => true, VERSIONS => 10}
```

Besides the default 'toStringBinary' format, 'scan' supports custom formatting by column. A user can define a `FORMATTER` by adding it to the column name in the scan specification. The `FORMATTER` can be stipulated:

1. either as a `org.apache.hadoop.hbase.util.Bytes` method name (e.g. `toInt`, `toString`)
2. or as a custom class followed by method name: e.g. `'c(MyFormatterClass).format'`.

Example formatting `cf:qualifier1` and `cf:qualifier2` both as Integers:

```
hbase> scan 't1', {COLUMNS => ['cf:qualifier1:toInt',
```



```
'cf:qualifier2:c(org.apache.hadoop.hbase.util.Bytes).toInt'] }
```

Note that you can specify a FORMATTER by column only (cf:qualifier). You cannot specify a FORMATTER for all columns of a column family.

Scan can also be used directly from a table, by first getting a reference to a table, like such:

```
hbase> t = get_table 't'
hbase> t.scan
```

Note in the above situation, you can still provide all the filtering, columns, options, etc as described above.

truncate

Disables, drops and recreates the specified table.

Examples:

```
hbase>truncate 't1'
```

4) HBase surgery tools

assign Assign a region. Use with caution. If region already assigned, this command will do a force reassign. For experts only.

Examples:

```
hbase> assign 'REGION_NAME'
```

balancer Trigger the cluster balancer. Returns true if balancer ran and was able to tell the region servers to unassign all the regions to balance (the re-assignment itself is async).

Otherwise false (Will not run if regions in transition).

Examples:

```
hbase> balancer
```

balance_switch Enable/Disable balancer. Returns previous balancer state.

Examples:

```
hbase> balance_switch true
```

```
hbase> balance_switch false
```

close_region Close a single region. Ask the master to close a region out on the cluster or if 'SERVER_NAME' is supplied, ask the designated hosting regionserver to close the region directly. Closing a region, the master expects 'REGIONNAME' to be a fully qualified region name. When asking the hosting regionserver to directly close a region, you pass the regions' encoded name only. A region name looks like

this:TestTable,0094429456,1289497600452.527db22f95c8a9e0116f0cc13c680396.The

trailing period is part of the regionserver name. A region's encoded name is the hash at the end of a region name; e.g. 527db22f95c8a9e0116f0cc13c680396 (without the period). A 'SERVER_NAME' is its host, port plus startcode. For example: host187.example.com,60020,1289493121758 (find servername in master ui or when you do detailed status in shell). This command will end up running close on the region hosting regionserver. The close is done without the master's involvement (It will not know of the close). Once closed, region will stay closed. Use assign to reopen/reassign. Use unassign or move to assign the region elsewhere on cluster. Use with caution. For experts only.

Examples:hbase> close_region 'REGIONNAME'

hbase> close_region 'REGIONNAME', 'SERVER_NAME'

compact Compact all regions in passed table or pass a region row to compact an individual region. You can also compact a single column family within a region.

Examples:

Compact all regions in a table:

hbase> compact 't1'

Compact an entire region:

hbase> compact 'r1'

Compact only a column family within a region:

hbase> compact 'r1', 'c1'

Compact a column family within a table:

hbase> compact 't1', 'c1'

flush Flush all regions in passed table or pass a region row to flush an individual region. For example:hbase> flush 'TABLENAME'

hbase> flush 'REGIONNAME'

major_compact Run major compaction on passed table or pass a region row to major compact an individual region. To compact a single column family within a region specify the region name followed by the column family name.

Examples:

Compact all regions in a table:

hbase> major_compact 't1'

Compact an entire region:

hbase> major_compact 'r1'

Compact a single column family within a region:

hbase> major_compact 'r1', 'c1'

Compact a single column family within a table:

hbase> major_compact 't1', 'c1'

move Move a region. Optionally specify target regionserver else we choose one

at random. NOTE: You pass the encoded region name, not the region name so this command is a little different to the others. The encoded region name is the hash suffix on region names: e.g. if the region name were TestTable,0094429456,1289497600452.527db22f95c8a9e0116f0cc13c680396. then the encoded region name portion is 527db22f95c8a9e0116f0cc13c680396

A server name is its host, port plus startcode. For example:

host187.example.com,60020,1289493121758

Examples:hbase> move 'ENCODED_REGIONNAME'

hbase> move 'ENCODED_REGIONNAME', 'SERVER_NAME'

split Split entire table or pass a region to split individual region. With the second parameter, you can specify an explicit split key for the region.

Examples:

split 'tableName'

split 'regionName' # format: 'tableName,startKey,id'

split 'tableName', 'splitKey'

split 'regionName', 'splitKey'

unassign Unassign a region. Unassign will close region in current location and then reopen it again. Pass 'true' to force the unassignment ('force' will clear all in-memory state in master before the reassign. If results in double assignment use hbck -fix to resolve. To be used by experts).

Use with caution. For expert use only. Examples:hbase> unassign 'REGIONNAME'

hbase> unassign 'REGIONNAME', true

hlog_roll Roll the log writer. That is, start writing log messages to a new file.

The name of the regionserver should be given as the parameter. A

'server_name' is the host, port plus startcode of a regionserver. For

example: host187.example.com,60020,1289493121758 (find servername in master ui or when you do detailed status in shell)

hbase>hlog_roll

zk_dump Dump status of HBase cluster as seen by ZooKeeper. Example:

hbase>zk_dump

5) Cluster replication tools

add_peer Add a peer cluster to replicate to, the id must be a short and the cluster key is composed like this:

hbase.zookeeper.quorum:hbase.zookeeper.property.clientPort:zookeeper.znode.parent

This gives a full path for HBase to connect to another cluster.

Examples:hbase> add_peer '1', "server1.cie.com:2181:/hbase"

hbase> add_peer '2', "zk1,zk2,zk3:2182:/hbase-prod"

remove_peer Stops the specified replication stream and deletes all the meta

information kept about it. Examples:

```
hbase> remove_peer '1'
```

`list_peers` List all replication peer clusters.

```
hbase> list_peers
```

`enable_peer` Restarts the replication to the specified peer cluster, continuing from where it was disabled.Examples:

```
hbase> enable_peer '1'
```

`disable_peer` Stops the replication stream to the specified cluster, but still keeps track of new edits to replicate.Examples:

```
hbase> disable_peer '1'
```

`start_replication` Restarts all the replication features. The state in which each stream starts in is undetermined.

WARNING:

start/stop replication is only meant to be used in critical load situations.

Examples:

```
hbase> start_replication
```

`stop_replication` Stops all the replication features. The state in which each stream stops in is undetermined.

WARNING:

start/stop replication is only meant to be used in critical load situations.

Examples:

```
hbase> stop_replication
```

6) Security tools

`grant` Grant users specific rights.

Syntax : `grantpermissions` is either zero or more letters from the set "RWXCA".

`READ('R'), WRITE('W'), EXEC('X'), CREATE('C'), ADMIN('A')`For example:`hbase> grant 'bobsmith', 'RWXCA'`

```
hbase> grant 'bobsmith', 'RW', 't1', 'f1', 'col1'
```

`revoke` Revoke a user's access rights.

Syntax : `revoke`

For example:

```
hbase> revoke 'bobsmith', 't1', 'f1', 'col1'
```

`user_permission` Show all permissions for the particular user.

Syntax : `user_permission`

For example:`hbase> user_permission`

```
hbase> user_permission 'table1'
```